

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“Jnana Sangama”, Belgaum– 590 014



**A Project Report
On**

**“ ENHANCING ACCIDENT PREDICTION THROUGH INTEGRATED
KNN AND DBSCAN ALGORITHMS FOR SUPERIOR ACCURACY ”**

Submitted in partial fulfillment of the requirements for the award of the degree of

**BACHELOR OF ENGINEERING
in
INFORMATION SCIENCE AND ENGINEERING**

Submitted by

ANKITHA K V

1KI22IS004

PREETHI KUMARI

1KI22IS037

LIKHITHA M N

1KI23IS401

Under the Guidance of

Prof. Chethan Kumar T B.E, M.Tech

Assistant Professor, Dept. of ISE

KIT, Tiptur



DEPARTMENT OF INFORMATION SCIENCE AND ENGINEERING

Kalpataru Institute of Technology, Tiptur– 572 201

2025-2026

Kalpataru Institute of Technology, Tiptur – 572 201

DEPARTMENT OF INFORMATION SCIENCE AND ENGINEERING



CERTIFICATE

This is to certify that the project work entitled “**ENHANCING ACCIDENT PREDICTION THROUGH INTEGRATED KNN AND DBSCAN ALGORITHMS FOR SUPERIOR ACCURACY**” has been successfully carried out by **ANKITHA K V (1KI22IS004)**, **PREETHI KUMARI (1KI22IS037)** and **LIKHITHA M N (1KI23IS401)** in partial fulfillment for the award of the degree of **Bachelor of Engineering in Information Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the academic year 2025-2026. It is certified that all corrections and suggestions indicated for internal assessment have been incorporated in the report. The project report has been approved as it satisfies the academic requirements with respect to the project work prescribed for the Bachelor of Engineering degree in Information Science and Engineering.

Guide

Prof. Chethan Kumar T BE, M.Tech
Assistant Professor
Dept. of ISE, KIT

HOD

Dr. Rajashekar K J M.Tech, Ph.D
Professor & Head
Dept. of ISE, KIT

Principal

Dr. Sateesh Kumar H C M.Tech, Ph.D
Principal, KIT

EXTERNAL VIVA VOCE

Name of the Examiners

Signature with Date

1. _____

2. _____

ACKNOWLEDGEMENT

The satisfaction and great happiness that accompany the successful completion of any task would be incomplete without mentioning about the people who made it possible. Here we make an honest attempt to express our deepest gratitude to all those who have been helpful and responsible for the successful completion of our project.

We would like to extend our heartfelt thanks to **Dr. H C Sathish Kumar, Principal**, Kalpataru Institute of Technology, Tiptur for his continuous support and encouragement throughout the course of this project.

We also express our sincere gratitude to **Dr. Rajashekar K J, Professor** and **HOD**, Department of Information Science and Engineering, Kalpataru Institute of Technology, Tiptur for his continuous guidance, valuable suggestion & encouragement during the course of our project.

We are deeply indebted to our internal guide, **Prof. Chethan Kumar T, Assistant Professor**, Department of Information Science and Engineering, Kalpataru Institute of Technology, Tiptur for his constant support, technical assistance and constructive suggestions and guidance for successful completion of our project. We are especially thankful to him for the motivation & enthusiasm that inspired us to complete our work successfully.

We would like to thank all faculty members of Department of Information Science and Engineering, KIT, Tiptur, our family members and our friends who have directly or indirectly contributed to our success.

Thanking You

ANKITHA K V 1KI22IS004
PREETHI KUMARI 1KI22IS037
LIKHITHA M N 1KI23IS401

ABSTRACT

The proposed system Enhancing Accident Prediction through Integrated KNN and DBSCAN Algorithms for Superior Accuracy aims to improve road safety by delivering a highly reliable and data-driven accident prediction framework. The model combines the strengths of the K- Nearest Neighbors (KNN) classifier with the Density-Based Spatial Clustering of Applications with Noise (DBSCAN) algorithm to effectively handle noisy, imbalanced, and heterogeneous traffic datasets. DBSCAN is utilized as a preprocessing mechanism to detect and eliminate outliers, ensuring that only meaningful and structurally consistent data contributes to the classification process. This refined dataset significantly enhances KNN's predictive capability by reducing misclassification caused by anomalous traffic patterns. The system processes multiple parameters, including vehicle speed, weather conditions, traffic density, visibility, and historical accident trends to generate accurate severity predictions. Through this hybrid integration, the model achieves superior accuracy, stability, and robustness compared to standalone traditional classifiers. The proposed framework demonstrates the potential of combining clustering and classification techniques to support intelligent transportation systems, enabling timely decision- making and contributing to accident prevention in real-world environments.

TABLE OF CONTENTS

CHAPTERS	PAGE NO
1 INTRODUCTION	1-5
1.1 Background of the study	2
1.2 Problem statement	3
1.3 Advantages and proposed system	4-5
2 LITERATURE SURVEY	6-8
3 SYSTEM REQUIREMENTS AND SPECIFICATION	9-18
3.1 Software requirement specification	9-11
3.2 Hardware requirement specification	12-13
3.3 Functional Requirements	13-15
3.4 Non-functional requirements	15-16
3.5 Data requirements	16-18
4 SYSTEM ANALYSIS	19-25
4.1 Problem identification	19-21
4.2 Objectives	22-23
4.3 Methodology	23-25
5 SYSTEM DESIGN	26-32
5.1 Architecture design	26-28
5.2 Detailed design	29-30
5.3 Output design	30-32
6 IMPLMENTATION	33-37
7 MODULE SPLIT	38-40
8 RESULTS AND SCREENSHOT	41-46
9 CONCLUSION AND FUTURE WORK	47-49
REFERENCES	50
BIBLIOGRAPHY	51
ANNEXURE	52

LIST OFFIGURES

FIGURE NO	TITLE	PAGE NO
5.1	Flowchart of System Design	28
8.1	Feature correlation heatmap	42
8.2	Accident severity distribution	43
8.3	Model confusion matrix	44
8.4	Accuracy comparision	45

Chapter 1

INTRODUCTION

Road accidents remain one of the most critical global challenges, causing severe injuries, fatalities, and economic losses each year. According to the World Health Organization (WHO), more than 1.3 million people die annually due to road traffic accidents, with millions more suffering long-term disabilities. These alarming statistics underscore the urgent need for effective accident prediction and prevention systems that can assist traffic authorities in mitigating risks and improving road safety. With the rapid development of data collection technologies such as sensors, GPS devices, and intelligent transportation systems, large volumes of traffic and environmental data are now available. This abundance of data provides an opportunity to apply advanced machine learning techniques to analyze patterns, identify accident-prone zones, and predict potential accidents before they occur.

Traditional statistical models and single-machine-learning algorithms have demonstrated limited performance in accident prediction due to challenges such as data imbalance, noise, and non-linear relationships among features. For example, the K-Nearest Neighbors (KNN) algorithm, while simple and effective for classification, is sensitive to noise and irrelevant data points. On the other hand, clustering algorithms such as DBSCAN (Density-Based Spatial Clustering of Applications with Noise) are highly capable of identifying spatial patterns and filtering noise from complex datasets. However, DBSCAN alone does not provide predictive capability. These limitations suggest that integrating complementary algorithms could yield more robust and accurate results.

This research proposes an integrated KNN–DBSCAN framework designed to enhance accident prediction accuracy by combining the strengths of both algorithms. In the proposed model, DBSCAN is first utilized to detect dense regions representing accident-prone areas and to remove outliers or noisy data points that may distort prediction results. Once the data is refined and clustered, the KNN algorithm is applied to classify accident likelihood based on the spatial and temporal characteristics of each cluster. This hybrid framework enhances prediction accuracy and robustness by effectively combining unsupervised clustering with supervised classification on complex accident datasets. The proposed approach supports proactive traffic risk assessment and informed decision-making for improved road safety management.

1.1 Background of the study

Traffic accidents have long been a pressing concern for both developed and developing nations, posing serious threats to human life, property, and the economy. The increasing rate of road accidents is often attributed to factors such as rapid urbanization, poor road conditions, excessive traffic volume, weather variations, and human error. According to global traffic safety reports, a substantial proportion of accidents occur in specific high-risk areas, commonly referred to as accident-prone zones or black spots. Identifying and predicting these zones in advance is crucial for implementing preventive measures. However, due to the dynamic and complex nature of traffic systems, accurately predicting accidents remains a challenging task that requires robust analytical models capable of handling large and diverse datasets. Traditional accident analysis methods have primarily relied on statistical and regression-based models, such as Poisson regression, logistic regression, and negative binomial models. While these approaches provide interpretability, they often struggle with nonlinear relationships and high-dimensional data, limiting their ability to capture the complex interactions among various factors that contribute to road accidents. In contrast, machine learning (ML) and data mining techniques have emerged as powerful alternatives, offering the ability to learn patterns from historical data and make data-driven predictions with higher precision. Algorithms such as Decision Trees, Random Forests, Support Vector Machines (SVM), and Neural Networks have been widely adopted in accident prediction studies. However, many of these methods still face challenges related to data noise, imbalance, and spatial heterogeneity.

Clustering and classification are two essential techniques in ML that can be effectively combined for accident prediction. Clustering algorithms, such as DBSCAN (Density-Based Spatial Clustering of Applications with Noise), group data points into clusters based on density, enabling the identification of accident hotspots and outliers. DBSCAN's ability to handle arbitrary-shaped clusters and detect noise makes it particularly suitable for traffic accident datasets, which often contain irregular patterns. However, clustering alone cannot predict future accident occurrences, highlighting the need for an integrated predictive framework. By combining DBSCAN with a classification technique such as KNN, accident-prone patterns can be identified more accurately while enabling effective severity prediction and risk assessment.

1.2 PROBLEM STATEMENT

Road traffic accidents represent a major global concern, resulting in significant loss of life, injuries, and economic damage each year. Despite numerous initiatives to enhance road safety, the frequency and severity of accidents continue to rise, particularly in rapidly urbanizing regions where vehicle density and human mobility are increasing. One of the major challenges in reducing accident rates lies in the inability to accurately predict when and where accidents are likely to occur. Accurate accident prediction is essential for implementing preventive measures, optimizing emergency response systems, and formulating effective traffic management strategies. However, the complexity of traffic environments, coupled with diverse influencing factors such as road geometry, weather, lighting conditions, and driver behavior, makes accident prediction a difficult task.

Existing predictive models, particularly traditional statistical and regression-based approaches, often fail to capture the nonlinear and dynamic relationships among variables involved in traffic accidents. These models typically assume a fixed data distribution and are sensitive to noise, outliers, and data imbalance—issues that are common in real-world accident datasets. Consequently, their predictive performance and generalization capabilities are limited. Machine learning algorithms have shown promise in addressing some of these limitations by learning complex patterns from data; however, single-model approaches often suffer from reduced accuracy and stability when confronted with noisy, sparse, or overlapping data points.

The K-Nearest Neighbors (KNN) algorithm is a widely used classification technique known for its simplicity and effectiveness in pattern recognition. Nonetheless, KNN's performance is highly dependent on data quality and is negatively affected by noise and irrelevant instances. On the other hand, DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a powerful clustering algorithm capable of identifying accident-prone regions and eliminating noise without requiring prior knowledge of the number of clusters. These limitations highlight the need for an integrated framework that combines clustering and classification to improve prediction reliability. Therefore, this project addresses the problem by proposing a hybrid DBSCAN–KNN model that enhances accident prediction accuracy by effectively handling noise, data imbalance, and complex spatial patterns.

1.3 ADVANTAGES AND PROPOSED SYSTEM

The integration of K-Nearest Neighbors (KNN) and DBSCAN (Density-Based Spatial Clustering of Applications with Noise) introduces several significant advantages over traditional and single-algorithm prediction models. One of the primary benefits is enhanced accuracy. By combining DBSCAN's ability to handle complex data distributions and noise with KNN's classification strength, the hybrid model minimizes the impact of irrelevant or misleading data points that often reduce predictive performance. DBSCAN effectively isolates clusters representing accident-prone zones and filters out noise, ensuring that only meaningful data are passed to the classification stage. This pre-processing step improves KNN's accuracy and reduces computational overhead.

Another major advantage is the robustness and adaptability of the model. Real-world traffic datasets are often unstructured, contain missing values, and exhibit non-linear relationships among features such as road type, weather, and traffic volume. The hybrid KNN-DBSCAN model can adapt to these variations by identifying dense clusters of accident occurrences and classifying them without relying on predefined parameters such as the number of clusters or distribution assumptions. This makes the model suitable for dynamic and heterogeneous traffic environments. Additionally, the integration improves noise tolerance, as DBSCAN's clustering mechanism inherently excludes outliers, preventing them from distorting the classification process.

The proposed model also enhances interpretability and decision support. By visualizing DBSCAN-generated clusters, transportation authorities can easily identify high-risk areas or —hotspots, while KNN's classification output provides predictive insights about the likelihood of future accidents in specific regions or under particular conditions. This dual-level analysis supports evidence-based decision-making in traffic management, allowing authorities to allocate resources effectively and implement proactive safety measures. Furthermore, the hybrid system can be scaled for real-time applications, enabling integration with intelligent transportation systems (ITS) for continuous monitoring and prediction.

1.3.1 Proposed System

The proposed system aims to develop a hybrid machine learning framework that integrates DBSCAN clustering and KNN classification to predict accident likelihood with superior accuracy. The system operates in several structured phases:

Data collection and preprocessing

Traffic accident data are collected from reliable sources such as government transportation databases, police records, and road sensor networks. The dataset includes attributes like location coordinates, weather conditions, time of day, vehicle type, and road features. Data cleaning, normalization, and feature selection are performed to Prepare the data for analysis.

Clustering Using DBSCAN:

The DBSCAN algorithm is applied to detect dense clusters representing accident- prone zones and to identify outliers or noise points. This process reduces dataset irregularities and enhances data quality by focusing on meaningful patterns.

Classification Using KNN:

The refined dataset obtained from DBSCAN is passed to the KNN algorithm. KNN classifies accident likelihood based on the nearest neighbors' historical patterns, producing an accurate prediction of potential accident occurrences.

Evaluation and Validation:

The system's performance is evaluated using metrics such as accuracy, precision, recall, and F1-score. Comparative analyses with individual algorithms (KNN and DBSCAN) and other conventional models are conducted to demonstrate the superiority of the hybrid framework.

Visualization and Decision Support:

The final output includes visual representations of accident-prone areas and prediction results. These insights assist traffic authorities in implementing preventive measures, optimizing patrol deployment, and improving road design and safety regulations.

Chapter 2

LITERATURE SURVEY

The field of traffic accident prediction has evolved significantly over the past two decades, driven by advances in machine learning (ML), data mining, and intelligent transportation systems (ITS). Researchers have explored numerous algorithms and models to improve prediction accuracy, handle noisy datasets, and identify accident-prone areas. This section reviews key studies and methodologies relevant to the proposed hybrid KNN–DBSCAN framework, highlighting their contributions, limitations, and areas for improvement.

[1] Traditional Statistical and Regression-Based Models

Early research on accident prediction primarily relied on statistical models, such as Poisson regression, logistic regression, and negative binomial models, to estimate accident frequency and severity. For instance, Lord and Mannering (2010) emphasized the use of count data models to analyze crash frequency, identifying variables like traffic volume, road type, and weather conditions as significant predictors. Similarly, Persaud and Nguyen (1998) used logistic regression to model accident probabilities at signalized intersections.

Although these models provided interpretability and insights into factor relationships, they assumed linear dependencies and specific data distributions. In real-world datasets, accident data are typically nonlinear, heterogeneous, and imbalanced, which limits the accuracy of purely statistical approaches. These limitations motivated the shift toward data-driven and non-parametric machine learning methods capable of learning complex patterns from large datasets.

[2] Machine Learning Techniques for Accident Prediction

With the growth of big data and advanced computing, machine learning techniques have become widely adopted in traffic safety research. Various supervised algorithms, including Decision Trees (DT), Random Forest (RF), Support Vector Machines (SVM), and Artificial Neural Networks (ANN), have been used to predict accident occurrence, severity, and risk levels.

Zhang et al. (2018) utilized a Random Forest model to predict highway accident severity based on environmental and roadway variables. Their study achieved higher accuracy than traditional regression models, highlighting the potential of ensemble methods. Similarly, Kumar and Toshniwal (2016) compared Naïve Bayes, Decision

Tree, and SVM models for accident severity prediction using real-world data from the National Highway Authority of India (NHAI). While SVM achieved the best accuracy, it required significant parameter tuning and was sensitive to outliers.

Other studies applied Neural Networks (NN) for accident prediction. Abdel-Aty and Pande (2005) employed a feed-forward NN to predict real-time crash likelihood on urban expressways, achieving promising results in detecting high-risk segments. However, NN-based methods often require large training datasets and are computationally expensive, limiting their use in real-time or resource-constrained environments.

While these ML models achieved improvements over statistical approaches, they still struggled with noisy data, imbalanced class distribution, and spatial heterogeneity—issues common in traffic accident datasets. These challenges led researchers to explore clustering-based and hybrid machine learning techniques.

[3] Clustering Techniques in Accident Analysis

Clustering algorithms group similar data points together, helping to identify spatial patterns and accident hotspots. Among them, DBSCAN (Density-Based Spatial Clustering of Applications with Noise) has proven particularly effective due to its ability to detect clusters of arbitrary shapes and to identify outliers automatically. Ester et al. (1996), who introduced DBSCAN, demonstrated its strength in handling spatial data with noise—a feature critical for traffic datasets. Later, Ma and Kockelman (2006) used DBSCAN to identify high-risk road segments and accident-prone zones using geospatial accident data. Their results showed that DBSCAN could successfully isolate clusters without requiring a predefined number of clusters, unlike K-Means.

Another study by Cheng et al. (2017) applied DBSCAN for identifying accident hotspots in urban areas, integrating it with Geographic Information Systems (GIS). The method provided visual insights into spatial accident patterns, but it lacked predictive capability it could describe where accidents were frequent but could not forecast future occurrences.

[4] K-Nearest Neighbors (KNN) in Accident Prediction

The K-Nearest Neighbors (KNN) algorithm is one of the simplest yet effective classification methods used in accident prediction. It classifies new instances based on their proximity to previously labeled data points. Shankar and Mannering (2012) applied KNN to predict accident severity levels, demonstrating that its non-parametric nature allows it to model complex, nonlinear relationships between input variables.

Similarly, Xu et al. (2019) used a KNN model to predict accident likelihood based on road geometry and traffic flow data. The model achieved reasonable accuracy but suffered from reduced performance when data contained noise or irrelevant features. Since KNN's performance heavily depends on the choice of distance metric and the value of k , it can be unstable in high-dimensional or noisy datasets.

Despite these drawbacks, KNN remains a strong baseline classifier due to its simplicity, interpretability, and efficiency for small to medium-sized datasets. The integration of a noise-handling algorithm such as DBSCAN can significantly improve KNN's performance by filtering out misleading data points before classification.

[5] Hybrid Models and Integrated Approaches

To address the limitations of single-algorithm approaches, researchers have recently focused on hybrid models that combine clustering and classification techniques. Hybrid frameworks can leverage the complementary strengths of multiple algorithms to improve prediction robustness and accuracy.

For example, Ahmed et al. (2020) proposed a hybrid K-Means–SVM model for traffic accident prediction, where K-Means was used for clustering and SVM for classification. The model outperformed standalone classifiers by reducing data redundancy. Similarly, Lee et al. (2021) combined DBSCAN and Random Forest to identify and classify accident hotspots, achieving improved precision in high-density traffic regions. However, Random Forest's complexity increased computational time, making it less practical for real-time applications.

Chapter 3

SYSTEM REQUIREMENTS AND SPECIFICATION

3.1 Software Requirements Specification.

The purpose of this Software Requirements Specification (SRS) document is to define the functional and non-functional requirements of the proposed accident prediction system that integrates the K-Nearest Neighbors (KNN) and Density-Based Spatial Clustering of Applications with Noise (DBSCAN) algorithms. The goal of this system is to analyze historical accident data, identify accident-prone areas, and predict the likelihood of future accidents with improved accuracy and reliability. This document outlines the objectives, operating environment, design constraints, and expected outcomes to guide the development and implementation of the system

Operating System

The proposed accident prediction system is designed to operate on a wide range of modern computer operating systems to ensure flexibility, accessibility, and ease of use. The preferred operating system for development and implementation is Windows 10 or later, as it provides a stable environment for data analysis and supports most of the required software tools and libraries, including Python, scikit-learn, and data visualization frameworks. Windows offers excellent compatibility with popular development environments such as Jupyter Notebook, PyCharm, and Spyder, making it suitable for both researchers and developers.

In addition to Windows, the system can also be efficiently executed on Linux-based operating systems, such as Ubuntu 20.04 or newer, which are commonly used in academic and research settings due to their high stability, open-source nature, and strong support for Python-based machine learning libraries. Linux provides better performance in handling large datasets and complex computational tasks, particularly when the system is deployed on servers or used for large-scale model training. The Linux environment also simplifies the installation of dependencies and supports automation through shell scripting, which is beneficial for repeated experiments and system testing.

Programming language

The proposed accident prediction system is developed using the Python programming language, which is widely recognized for its simplicity, flexibility, and powerful support for data science and machine learning applications. Python has become the preferred language for artificial intelligence research due to its extensive

library ecosystem, strong community support, and ease of integration with analytical and visualization tools. Its clear and readable syntax enables developers and researchers to focus on algorithm design and data analysis rather than complex programming structures. Python provides a variety of specialized libraries that make it particularly suitable for implementing machine learning algorithms such as K-Nearest Neighbors (KNN) and DBSCAN (Density-Based Spatial Clustering of Applications with Noise). The scikit-learn library is used for implementing both algorithms, offering optimized and well-tested functions for clustering, classification, and model evaluation. NumPy and Pandas are utilized for data manipulation and numerical computation, enabling efficient handling of large and complex accident datasets. For visualization purposes, libraries such as Matplotlib and Seaborn are employed to generate cluster plots, prediction graphs, and performance charts, which assist in interpreting the results and identifying accident-prone areas.

Python also supports seamless integration with data storage systems and web-based applications, allowing for future scalability of the project. Through frameworks like Flask or Django, the developed model can later be deployed as a web service or integrated into real-time intelligent transportation systems (ITS). Additionally, Python's compatibility with multiple operating systems—Windows, Linux, and macOS—ensures that the system remains platform-independent and easily portable.

Another key reason for choosing Python is its strong suitability for research and prototyping. It allows for rapid experimentation with different parameters, algorithms, and data preprocessing techniques without requiring extensive recompilation or complex configurations. This flexibility is essential in an academic and experimental setting, where the goal is to test, compare, and optimize machine learning models for improved accuracy.

Libraries and Frameworks.

The development of the proposed accident prediction system relies heavily on a set of powerful Python libraries and frameworks that provide the essential tools for data preprocessing, machine learning, visualization, and performance evaluation. These libraries form the backbone of the system, enabling efficient implementation of the K-Nearest Neighbors (KNN) and DBSCAN (Density-Based Spatial Clustering of Applications with Noise) algorithms while ensuring scalability, accuracy, and ease of maintenance. unsupervised learning. Scikit-learn provides optimized functions for clustering, classification, model evaluation, and hyperparameter tuning, making it ideal

for implementing both KNN and DBSCAN. It also includes utilities for data preprocessing, such as normalization, feature scaling, and cross-validation, which are crucial for improving model performance and accuracy.

For data manipulation and analysis, the system utilizes Pandas and NumPy. Pandas is used to manage and process structured datasets, offering high-level data structures such as Data Frames that simplify handling of large accident datasets with multiple features. It allows for easy loading, cleaning, transformation, and merging of data from various sources. NumPy, on the other hand, provides support for numerical and matrix operations, allowing efficient computation of distances and similarity measures required by the KNN algorithm. Together, these libraries ensure that the system can efficiently manage large-scale datasets and perform computations with high accuracy.

Visualization plays a vital role in interpreting the results of the model and identifying accident-prone zones. The Matplotlib and Seaborn libraries are used to create graphical representations such as scatter plots, heatmaps, and cluster visualizations. These tools help users understand spatial relationships and the distribution of accident occurrences across different regions. Seaborn, built on top of Matplotlib, provides aesthetically pleasing and informative statistical visualizations that make it easier to present the results of clustering and classification in research or decision-making contexts.

In addition to these libraries, GeoPandas may be used for spatial data visualization and mapping of accident locations. GeoPandas extends the Pandas library to handle geographical data and allows easy plotting of accident clusters on real-world maps, providing a clear spatial understanding of accident hotspots.

For model evaluation and validation, Scikit-learn's metrics module is employed to calculate performance indicators such as accuracy, precision, recall, F1-score, and ROC-AUC. These metrics are essential for assessing the effectiveness of the hybrid model and comparing it against standalone algorithms.

The proposed system can also be extended using frameworks like Flask or Streamlit for web-based deployment and user interaction. Flask provides a lightweight web framework for integrating the trained model into a web application, enabling real-time accident prediction.

Streamlit, a Python-based dashboarding framework, can be used to build interactive visualization interfaces without extensive web development skills, making it ideal for

researchers and policymakers.

Furthermore, for experimental management and reproducibility, Jupyter Notebook serves as the main development and testing environment. It provides an interactive interface that supports code execution, documentation, and visualization within a single workspace, facilitating step-by-step experimentation with the hybrid KNN–DBSCAN model.

3.2 Hardware Requirement Specification.

The proposed hybrid accident prediction system has been designed to operate efficiently on standard computing hardware without the need for specialized or high-end equipment. However, to ensure smooth execution of data processing, clustering, and classification tasks, certain minimum and recommended hardware specifications are necessary. These requirements are intended to support both the development and testing phases of the system as well as potential deployment for larger datasets or research environments.

At a minimum, the system requires a personal computer or workstation equipped with a modern multi-core processor, such as an Intel Core i5 or an AMD Ryzen 5, capable of handling multiple computational threads simultaneously. Since both the K- Nearest Neighbors (KNN) and DBSCAN algorithms involve iterative calculations and matrix operations, having a processor with multiple cores significantly enhances the speed and efficiency of the computations. For optimal performance, a higher-end processor such as an Intel Core i7 or AMD Ryzen 7 is recommended, especially when dealing with large accident datasets or performing multiple experiments concurrently.

Adequate memory capacity (RAM) is essential for the smooth functioning of the system. A minimum of 8 gigabytes (GB) of RAM is required to handle moderate-sized datasets and ensure efficient execution of Python-based data analysis libraries. However, for advanced users or institutions processing large, high-dimensional datasets—especially those containing geographical or temporal features—16 GB of RAM or more is recommended. Sufficient memory allows the system to store data in memory for quick access, reducing lag during training and evaluation phases.

In terms of storage, at least 50 GB of available disk space is necessary for installing the required software libraries, storing datasets, and saving model outputs and visualizations. Since accident datasets may include multiple years of historical data or detailed records with numerous attributes, a solid-state drive (SSD) is recommended over a traditional hard disk drive (HDD) to ensure faster data retrieval and read/write speeds.

SSD storage also improves overall system responsiveness during preprocessing and visualization tasks.

While a dedicated graphics processing unit (GPU) is not a strict requirement for this project—since KNN and DBSCAN are primarily CPU-based algorithms—a GPU can be beneficial if the system is later extended to include deep learning or neural network components. A mid-range NVIDIA GPU would enhance computational speed in such advanced applications or when dealing with large-scale geospatial visualizations.

Additional hardware such as a high-resolution monitor is recommended to facilitate clear visualization of clusters, maps, and statistical graphs generated during analysis. A stable internet connection is also useful for downloading Python packages, accessing external datasets, or integrating with cloud-based platforms for data storage and computation.

The system should operate reliably on standard desktop or laptop configurations used in Academic and research institutions. For larger-scale implementations, such as those deployed by government agencies or transport authorities, the model can be hosted on dedicated servers or cloud platforms equipped with higher processing power, expanded memory, and greater storage capacity to handle continuous data inflow and real-time predictions.

In summary, the proposed accident prediction system does not require specialized hardware but benefits from moderate computational resources to achieve optimal performance. A system with a modern multi-core processor, at least 8 GB of RAM (preferably 16 GB), and SSD storage ensures efficient data handling and accurate model execution. These specifications make the project feasible and accessible to researchers, students, and organizations seeking to implement data-driven accident prediction and analysis using the hybrid KNN–DBSCAN framework.

3.3 Functional Requirements

The functional requirements define the specific operations and behaviors that the proposed accident prediction system must perform to achieve its intended objectives. These requirements describe how the system will process data, execute algorithms, generate outputs, and interact with users or other components. The system integrates DBSCAN for clustering and KNN for classification to predict accident likelihood with improved accuracy and reliability.

The primary function of the system is data acquisition and preprocessing. The

system must be capable of importing traffic accident data from various sources such as government databases, CSV or Excel files, or sensor data logs. Once imported, the data should undergo a preprocessing phase that includes cleaning missing or inconsistent entries, handling duplicate records, and normalizing numerical features to ensure uniform scaling. This process ensures that the dataset is properly formatted and suitable for machine learning analysis.

The next essential function is data clustering using the DBSCAN algorithm. The system should analyze the spatial and temporal characteristics of accident data to identify clusters of high accident density, known as accident-prone zones or hotspots. DBSCAN must automatically determine dense regions based on distance (epsilon) and minimum point thresholds (minPts), grouping closely related data points together while identifying outliers as noise. This function enables the system to isolate meaningful patterns and eliminate irrelevant data that may negatively impact model accuracy.

After clustering, the system must perform classification using the K-Nearest Neighbors (KNN) algorithm. The refined dataset produced by DBSCAN serves as input for the KNN model, which classifies new or unseen data points based on their similarity to previously observed accident records. KNN predicts the likelihood or severity of an accident by analyzing the features of the k nearest data points. This classification function allows the system to produce accurate predictions that can help identify potential accident occurrences in the future. Another key functional requirement is parameter tuning and optimization. The system should allow users to adjust key parameters such as the epsilon and minPts values for DBSCAN, and the k value for KNN. These parameters play a crucial role in model performance, and the system must enable experimentation to find optimal settings that maximize prediction accuracy and minimize false results. The system must also provide model evaluation and validation capabilities. After training and testing, the model's performance should be assessed using evaluation metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. The evaluation process helps determine the effectiveness of the hybrid model and compare its performance against standalone algorithms or other machine learning techniques.

A further important functionality is data visualization. The system should generate visual representations of results, including accident clusters, prediction graphs, and hotspot maps.

Visualization assists users in understanding spatial relationships, model

performance, and accident trends, making the results easier to interpret. Charts, scatter plots, and heatmaps created using visualization libraries should be clear, informative, and suitable for presentation or reporting. The system must also incorporate result storage and export functions. Once the model has generated predictions and evaluation metrics, these results should be stored in a structured format for future reference. The system should support exporting reports or visualizations in common file formats such as CSV, PDF, or PNG, enabling users to share findings with other stakeholders.

In addition, the system should offer user interaction and control through a simple and intuitive interface or notebook environment. Users should be able to input data, adjust algorithm parameters, run analyses, and view outputs without requiring extensive programming knowledge. The interface should guide users through each stage of the process, from data loading to visualization. Finally, the system must support future scalability and integration. Although initially designed for offline analysis, it should be extendable to integrate with real-time data sources such as IoT sensors, GPS devices, or traffic surveillance systems. This functionality would allow the system to provide continuous accident prediction and alerting capabilities, making it valuable for smart transportation and public safety applications.

3.4 Non-Functional Requirements.

The non-functional requirements define the overall qualities, performance standards, and operational constraints that the proposed accident prediction system must satisfy. Unlike functional requirements, which specify what the system does, non-functional requirements describe how efficiently, reliably, and securely the system performs its tasks. These requirements ensure that the system not only produces accurate predictions but also operates effectively across different environments and user scenarios.

One of the key non-functional aspects of the system is **performance**. The system must be capable of processing large volumes of accident data efficiently without causing delays or computational bottlenecks. Since both the DBSCAN and KNN algorithms involve complex calculations, the system should be optimized to handle datasets containing thousands or even hundreds of thousands of records within a reasonable timeframe. Efficient data handling and algorithmic optimization will ensure that the prediction process remains fast and responsive, even when executed on standard computing hardware.

Another crucial requirement is **accuracy**. The hybrid approach combining DBSCAN

and KNN is specifically designed to enhance predictive performance compared to standalone models. The system should achieve a high level of accuracy, precision, and recall when predicting accident likelihood, ideally exceeding traditional single-model approaches by at least 10%. Continuous testing and parameter tuning should be conducted to maintain and improve model performance over time. platform, secure data transmission protocols should be implemented to prevent unauthorized access or data breaches.

Another important non-functional requirement is **compatibility**. The system should be designed to operate across multiple operating systems, including Windows, Linux, and macOS, to maximize accessibility for different users. Compatibility with various versions of Python and commonly used machine learning libraries must also be ensured to prevent dependency issues and maintain smooth execution.

Efficiency in terms of resource utilization is equally significant. The system should be optimized to minimize memory usage and CPU load while maintaining high computational accuracy. This ensures that the model can be executed on standard hardware without requiring specialized equipment.

Finally, the system must also meet the requirement of **portability**. It should be easily deployable on different hardware configurations and adaptable to various computing environments. Whether used in academic research, government analysis, or intelligent transportation systems, the model should maintain consistent functionality and performance.

3.5 Data Requirements.

The effectiveness of the proposed hybrid accident prediction system depends heavily on the availability, quality, and completeness of the input data. Accurate predictions and meaningful clustering rely on comprehensive datasets that capture all relevant aspects of traffic accidents, including temporal, spatial, environmental, and vehicular factors. Therefore, the system has specific data requirements that must be satisfied to ensure reliable model performance.

The system requires historical accident data as the primary input. This data should ideally contain records of past accidents over multiple years to capture trends and patterns. Each record must include essential attributes such as the date and time of the accident, precise geographic location (latitude and longitude), type and severity of the accident, number and types of vehicles involved, and the number of casualties or injuries.

These features are critical for the DBSCAN algorithm to identify accident clusters and for the KNN algorithm to classify and predict the likelihood of future accidents.

In addition to temporal and spatial information, environmental and road conditions are important data attributes. This includes weather conditions (rain, fog, snow, or clear), lighting conditions (day, night, or dusk), road type (highway, city road, or rural road), traffic density, and road surface conditions. These factors influence accident occurrence and severity, and incorporating them enhances the predictive power of the model.

Data preprocessing is a key requirement to ensure quality and consistency. The system expects data to be free of missing, inconsistent, or duplicate entries, or it must be capable of handling such issues during preprocessing. Numerical features should be normalized or standardized as necessary, and categorical features should be encoded appropriately for use in machine learning algorithms. Proper preprocessing ensures that the KNN distance calculations and DBSCAN clustering are accurate and meaningful.

The system may also require external supplementary data for enhanced analysis. This could include traffic volume datasets, GPS-based vehicle movement records, and demographic data for accident-prone regions. Integrating these datasets allows for more precise clustering and more accurate classification results, improving the system's overall performance and reliability. In terms of format, the system supports data in common structured formats such as CSV, Excel, or SQL databases. Each dataset must be organized with clearly labeled columns corresponding to the relevant attributes. For geographic data, the inclusion of latitude and longitude coordinates is essential to facilitate spatial clustering by DBSCAN and accurate mapping of accident hotspots.

Furthermore, data volume and frequency are important considerations. While the system can operate on moderate-sized datasets for offline analysis, larger datasets are required for training robust models capable of handling real-world scenarios. If the system is extended for real-time prediction, it must be able to ingest streaming data from sensors, traffic cameras, or other sources with minimal latency.

Finally, data integrity and security are critical requirements. All datasets used must be reliable, accurate, and obtained from trusted sources such as government transport departments, traffic authorities, or verified accident databases. The system should also ensure secure storage and processing of sensitive information to prevent unauthorized access or data breaches.

In summary, the data requirements for the hybrid KNN–DBSCAN accident prediction system emphasize completeness, accuracy, consistency, and relevance. The

system requires detailed historical accident data enriched with environmental, vehicular, and traffic-related attributes. Proper preprocessing, secure handling, and integration with supplementary datasets ensure that the model can produce accurate predictions, identify accident-prone zones, and provide actionable insights for road safety planning and management.

3.6 Proposed System

The proposed accident prediction system is designed using a hybrid architecture that integrates DBSCAN clustering and K-Nearest Neighbors (KNN) classification to achieve superior prediction accuracy. The architecture follows a modular, step-by-step workflow that ensures efficient data processing, accurate analysis, and clear visualization of results. This design allows the system to handle large datasets, extract meaningful patterns, and generate actionable insights for traffic management and accident prevention. The system begins with a data acquisition module, where historical accident data and supplementary datasets are collected from multiple sources, including government traffic databases, sensors, and CSV or Excel files. This raw data contains temporal, spatial, environmental, and vehicular information, which forms the basis for predictive analysis. Data acquisition is followed by a **data** preprocessing module, which cleanses the data by removing duplicates, handling missing values, and normalizing numerical features. Categorical attributes, such as road type or weather conditions, are encoded into numerical values to ensure compatibility with machine learning algorithms. This preprocessing step is crucial to improve the accuracy of both clustering and classification models.

Chapter 4

SYSTEM ANALYSIS

4.1 Problem Identification

Road accidents are a major concern worldwide, leading to loss of life, injuries, property damage, and economic costs. Despite the availability of traffic regulations, surveillance systems, and preventive measures, accidents continue to occur frequently due to a combination of human, environmental, and infrastructural factors. One of the primary challenges in addressing road safety is the lack of effective predictive models that can accurately identify accident-prone areas and anticipate potential accident occurrences. Traditional methods often rely on simple statistical analysis or historical records without leveraging advanced machine learning techniques, resulting in limited predictive accuracy and delayed interventions.

Another problem is the complexity and heterogeneity of accident data. Traffic accident datasets often include spatial coordinates, timestamps, vehicle information, weather conditions, road type, and traffic density. These multidimensional features make it difficult for conventional algorithms to detect meaningful patterns. Moreover, the presence of noise, outliers, and incomplete data further reduces the effectiveness of predictive models. Existing systems frequently fail to distinguish between high-risk clusters and sporadic incidents, leading to inaccurate predictions and inefficient resource allocation for accident prevention.

Additionally, many current approaches focus on either unsupervised learning for clustering accident-prone zones or supervised learning for classification separately, without integrating the strengths of both methodologies. Clustering methods alone may identify regions with high accident density but cannot predict the likelihood of future accidents. Conversely, classification models can predict outcomes based on historical data but may be influenced by outliers and noisy data points, resulting in lower accuracy. The lack of a hybrid approach that combines clustering and classification limits the practical utility of existing accident prediction systems.

The problem is further compounded by the need for real-time or near real-time predictions, which is increasingly important for traffic management authorities emergency services, and smart city initiatives. Existing models often require extensive preprocessing and manual intervention, making them unsuitable for dynamic traffic environments. Inadequate visualization and reporting mechanisms also hinder decision-

makers from understanding trends, identifying hotspots, and implementing timely preventive measures.

4.1.1 Challenges

Accident prediction is a highly complex problem due to the nature and variability of traffic and environmental data. One of the primary challenges is the high dimensionality of accident data. Accident records typically include multiple attributes such as time, date, geographic location, vehicle types, road conditions, weather, traffic density, and accident severity. Managing and analyzing such multidimensional data requires sophisticated algorithms capable of extracting meaningful patterns without being overwhelmed by irrelevant or redundant information.

Data quality and completeness represent another significant challenge. Real-world accident datasets often have missing entries, inconsistent values, duplicate records, or even erroneous data due to human reporting errors or sensor malfunctions. If these data issues are not addressed through preprocessing, they can negatively impact both the clustering and classification processes, leading to inaccurate predictions and misidentification of accident-prone areas.

The presence of noise and outliers is a further complication. Rare or unusual accident events may not conform to the typical patterns in the data, potentially distorting the clustering results of DBSCAN or the neighbor-based classification of KNN. Accurately distinguishing between meaningful patterns and random anomalies is essential to prevent misleading predictions.

Algorithm integration also poses a challenge. DBSCAN, an unsupervised clustering algorithm, and KNN, a supervised classification algorithm, have inherently different operating mechanisms. Ensuring that the clusters generated by DBSCAN effectively feed into KNN for accurate prediction requires careful system design, proper data preprocessing, and consistent feature representation. Misalignment between these stages can reduce the overall system performance. Parameter optimization is another crucial issue. DBSCAN requires the selection of epsilon (radius for neighborhood) and minPts (minimum points to form a cluster), while KNN requires the k-value (number of neighbors) and distance metrics. Improper selection of these parameters can result in poor clustering, misclassification, or overfitting, making it essential to implement methods for systematic tuning and validation. Scalability becomes a challenge when dealing with large datasets.

Running DBSCAN and KNN on such large datasets without optimizing for computational efficiency may lead to long processing times, excessive memory usage, or even system crashes. Real-time prediction requirements introduce additional complexity. For traffic authorities and emergency services, predicting accidents in near real-time is valuable. However, processing incoming traffic sensor data, GPS coordinates, and environmental information quickly while performing clustering and classification simultaneously demands efficient algorithms and resource management.

Another challenge is visualization and interpretability. Even with accurate predictions, results must be presented in a clear and actionable way. Accident-prone zones, hotspot maps, and predictive risk charts need to be understandable to traffic authorities, city planners, and policymakers who may not have technical expertise in machine learning.

Heterogeneous data integration is also challenging. Accident prediction often requires combining data from multiple sources such as GPS sensors, weather APIs, traffic cameras, and historical accident databases. These datasets may have different formats, units, or temporal resolutions, requiring careful preprocessing and synchronization before they can be used effectively.

Ensuring model reliability and consistency is critical. The system must produce reproducible results, providing confidence that predictions are trustworthy. Inconsistent outputs or poor handling of new, unseen data can undermine user trust and reduce the practical utility of the system. Handling imbalanced datasets is another important challenge. Severe or fatal accidents occur less frequently than minor incidents, making it difficult for KNN to classify rare events accurately without bias toward more common cases. Techniques such as oversampling, weighting, or synthetic data generation may be required to address this issue.

Environmental factors, such as sudden weather changes or road construction, introduce unpredictability, complicating accurate prediction. Furthermore, hardware and resource limitations can hinder the ability to process large volumes of data efficiently, especially when real-time predictions are required. Finally, security, data privacy, and adaptability pose additional challenges. Sensitive accident data must be securely stored and processed to prevent unauthorized access, and the system should be adaptable to different geographical regions, road infrastructures, and traffic behaviors without extensive reconfiguration.

4.2 OBJECTIVES

The primary objective of this project is to develop a robust and accurate accident prediction system that integrates the strengths of both DBSCAN clustering and K-Nearest Neighbors (KNN) classification. By combining these algorithms, the system aims to overcome the limitations of traditional methods that either rely solely on clustering or classification, thereby enhancing predictive accuracy and reliability. A key goal is to identify accident-prone areas effectively and predict potential future accidents based on historical patterns and relevant environmental factors.

A core objective is to handle complex, multidimensional accident datasets efficiently. The system is designed to preprocess and clean data, manage missing or inconsistent entries, and normalize numerical attributes while encoding categorical variables, ensuring that both clustering and classification algorithms can operate effectively. This objective directly addresses challenges related to data quality, noise, and heterogeneity.

Another important objective is to detect and manage outliers and anomalies in accident data. By applying DBSCAN clustering first, the system can separate dense accident clusters from sparse or isolated points, allowing KNN to focus on meaningful patterns and reduce the influence of rare or erroneous data entries. This approach enhances model robustness and reduces false predictions. The project also aims to optimize algorithmic parameters to maximize system performance. For DBSCAN, this involves selecting the most suitable epsilon and minPts values, while for KNN, determining the optimal k-value and distance metric. Systematic parameter tuning ensures that the hybrid model achieves the best balance between accurate clustering and classification.

Scalability and computational efficiency are further objectives. The system is designed to process large-scale datasets efficiently without compromising speed or accuracy, making it suitable for city-wide or national traffic datasets. Relatedly, real-time or near real-time prediction capabilities are targeted, enabling traffic authorities and emergency services to respond proactively to potential accident risks.

Another key objective is to provide intuitive visualization and actionable insights. The system aims to generate clear graphical representations, including accident heatmaps, cluster plots, and risk prediction charts, to support decision-making by traffic planners and policymakers.

The project also seeks to ensure system adaptability and reliability. The hybrid model should perform consistently across different regions, traffic conditions, and accident types while , maintaining secure handling of sensitive data. Additionally, it aims to address challenges such as imbalanced datasets, rare severe accidents, and environmental variability, making the system robust in real-world scenarios.

In summary, the project objectives are to:

1. Develop a hybrid DBSCAN–KNN accident prediction system with superior accuracy.
2. Efficiently handle high-dimensional, heterogeneous, and noisy accident data.
3. Detect accident-prone clusters and predict future accident likelihood.
4. Optimize algorithm parameters for improved clustering and classification performance.
5. Ensure scalability and near real-time prediction capabilities.

4.3 Methodology

The methodology of the proposed system is designed to build a hybrid accident prediction framework that integrates DBSCAN clustering and K-Nearest Neighbors (KNN) classification, ensuring superior accuracy, robustness, and actionable insights. The process is divided into multiple stages, each addressing specific challenges identified in accident prediction.

Data Acquisition and Integration:

The first stage involves collecting comprehensive traffic accident datasets from multiple sources, including government accident records, police reports, GPS- enabled vehicles, traffic sensors, and weather monitoring systems. The collected data typically includes features such as date, time, location coordinates, road type, vehicle type, traffic density, weather conditions, and accident severity. This stage also involves data integration, where heterogeneous sources are merged into a single dataset, resolving inconsistencies in formats, units, and timestamps to ensure uniformity.

Data Preprocessing:

Once collected, the raw data undergoes rigorous preprocessing to improve quality and reliability. This includes data cleaning to remove duplicates, correct inconsistencies, and handle missing values using imputation techniques. Feature engineering is applied to extract relevant information, such as converting timestamps into time-of-day

categories. calculating traffic congestion levels from raw sensor data. Normalization and encoding are performed to scale numerical data and convert categorical features into machine-readable formats. Noise reduction techniques are applied to minimize the effect of outliers on subsequent clustering and classification.

Clustering with DBSCAN:

The refined dataset is first analyzed using DBSCAN (Density-Based Spatial Clustering of Applications with Noise), an unsupervised algorithm that groups accident records into dense clusters while marking sparse points as noise. This step identifies accident-prone areas, enabling focused analysis on regions with high accident density. Proper selection of DBSCAN parameters—epsilon (ϵ) and minimum points (minPts)—is crucial for accurate clustering. DBSCAN's ability to handle irregular cluster shapes and noisy data makes it suitable for spatial accident data, which often does not conform to simple patterns.

Classification with KNN:

Following clustering, the system applies K-Nearest Neighbors (KNN) for predictive modeling. KNN predicts the probability or severity of accidents based on the similarity of new input instances to existing patterns in the clustered dataset. The number of neighbors (k) and the distance metric (e.g., Euclidean or Manhattan) are optimized to achieve the highest classification accuracy. By combining DBSCAN and KNN, the system benefits from unsupervised pattern discovery and supervised predictive accuracy, allowing it to identify both high-risk areas and the likelihood of future accidents.

Handling Imbalanced and Rare Events:

Accident datasets often exhibit class imbalance, with severe accidents being rarer than minor ones. To address this, the methodology incorporates techniques such as oversampling, undersampling, or synthetic data generation (SMOTE) to ensure the KNN classifier does not become biased toward majority classes. This enhances the system's ability to predict rare but critical incidents.

Model Evaluation and Validation:

The hybrid model is rigorously evaluated using performance metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. Cross-validation is applied to verify robustness and prevent overfitting. The evaluation process ensures that the system can generalize well to new and unseen data, maintaining reliability in real-world deployment.

Visualization and Reporting:

The final stage involves data visualization to make the results actionable for traffic authorities and decision-makers. Accident clusters are displayed using heatmaps, risk zones are highlighted on maps, and prediction results are summarized in charts and tables.

Chapter 5

SYSTEM DESIGN

5.1 Architectural Design

The architectural design of the proposed accident prediction system is structured as a multi-layered framework that integrates data acquisition, preprocessing, hybrid machine-learning analytics, and alert generation to achieve superior prediction accuracy and real-time decision support. Each module in the architecture works collaboratively to ensure smooth data flow, high computational efficiency, and improved reliability when combining the strengths of KNN and DBSCAN algorithms.

The first layer, the **City & Vehicle Data Collection Unit**, gathers heterogeneous data from various sources such as city traffic sensors, GPS-enabled vehicles, environmental monitoring systems, and historical accident datasets. This unit collects critical features including temperature, pressure, location coordinates, real-time traffic movement, and past accident records. By aggregating multi-dimensional and multi-source inputs, this layer establishes a comprehensive and reliable dataset essential for building accurate accident prediction models.

The second layer, the **Pre-processing Module**, transforms raw data into a clean and standardized format suitable for machine-learning operations. This involves scaling, normalization, noise removal, and filtering inconsistent or missing values. The module ensures that the data fed into the prediction engine is uniform, noise-free, and properly scaled, allowing distance-based algorithms like KNN and DBSCAN to operate effectively without bias caused by varying data magnitudes.

The third layer, the **Accident Prediction Engine (KNN Algorithm)**, performs supervised classification to predict accident likelihood or severity. KNN analyzes the input features by comparing them with historical patterns using optimized k-values and distance metrics. The model identifies similarities in conditions such as traffic density, road environment, and vehicle behavior, generating accurate predictions based on nearest neighbors. This layer outputs classification results, cluster tendencies, and noise indicators that assist in further analysis.

The fourth layer, the **Accident Clustering Engine (DBSCAN Algorithm)**, applies density-based clustering to detect accident-prone zones and isolate anomalous cases. DBSCAN evaluates the distribution of accident points using ϵ (epsilon) distance and minimum sample thresholds to form meaningful clusters while marking low-density

or irregular points as noise. This step reduces the influence of outliers on the final decision-making process and contributes to improved stability and reliability of the model.

The fifth layer, the **Integrated Engine**, fuses outputs from both KNN and DBSCAN to produce a consolidated risk assessment. By combining KNN's classification accuracy with DBSCAN's noise filtering and hotspot detection, this engine produces enhanced prediction results, performance metrics, and integrated risk scores. This hybrid integration strengthens decision-making by reducing false positives and improving sensitivity toward high-risk patterns.

The final operational layer, the **Alert Module**, transforms the integrated risk scores into actionable outputs. It generates alerts indicating accident probability or severity levels, enabling real-time notifications for drivers, traffic authorities, or intelligent transportation systems. These alerts help in preventive action, traffic routing, and situational awareness, ultimately contributing to road safety enhancement.

Overall, the proposed architectural design follows a structured, modular approach that maps precisely to the system diagram. By seamlessly integrating data acquisition, preprocessing, clustering, classification, and alert generation, the framework ensures scalability, adaptability, and high predictive performance. The combination of DBSCAN and KNN within this layered architecture enables robust accident prediction and provides meaningful, actionable insights for intelligent traffic management and public safety.

Furthermore, the integrated architecture is designed to leverage the complementary strengths of both supervised and unsupervised learning. DBSCAN plays a critical role in identifying structural patterns within the data by discovering dense accident hotspots and isolating noise points that could negatively influence model performance. This significantly enhances the quality of the dataset that flows into the KNN classifier. The KNN model, being distance-based, benefits greatly from the outlier removal performed by DBSCAN because it ensures that the nearest neighbors contributing to classification are genuinely representative of real accident patterns. This hybrid interaction not only improves prediction accuracy but also enhances the overall stability and generalizability of the system in diverse road environments.

In addition, the architecture supports continuous learning through iterative updates. As new accident records, sensor values, or environmental conditions are added to the system, the preprocessing, clustering, and classification models can be retrained or

updated to maintain high accuracy over time. This makes the system suitable for real- world deployment where traffic patterns and road conditions change dynamically. The modularity of the system also enables easy integration of additional machine-learning models or data sources without disrupting existing components, ensuring long-term scalability and adaptability.

Another significant advantage of this design is its capability for real-time processing. The efficient flow of data from collection to alert generation allows the system to operate at low latency, making it capable of issuing timely warnings that could potentially prevent accidents. The alert module functions as the final actionable layer, converting prediction results into risk notifications that can assist drivers, traffic control centers, and smart city infrastructures in taking preventive measures. This emphasizes the practical applicability of the system beyond theoretical prediction, making it a valuable tool for enhancing road safety. Overall, this extended architecture not only supports accurate accident prediction but also ensures practical usability, adaptability, and continuous improvement. By integrating DBSCAN and KNN within a well-structured layered design, the system provides a robust solution capable of meeting the evolving needs of modern intelligent transportation systems.




Fig 5.1 Flowchart of System Design

5.2 DETAILED DESIGN

The detailed design of the proposed accident prediction system builds upon the architectural framework, specifying how each component functions, interacts, and contributes to achieving accurate predictions. The system is divided into multiple modules, each responsible for a distinct task, and collectively forming a robust pipeline from raw data to actionable insights.

Data Acquisition Module

This module is responsible for gathering accident-related data from multiple sources, including government traffic databases, police records, GPS-enabled vehicles, traffic sensors, and weather monitoring systems. The module handles heterogeneous data integration, ensuring that records from different sources are compatible in terms of format, timestamp, and measurement units. For instance, GPS coordinates are standardized, timestamps are converted to a consistent time zone, and weather data is matched to accident times. The module also supports automated data updates, enabling continuous ingestion of new records.

Data Preprocessing Module

The preprocessing module prepares raw data for clustering and classification. Key functions include data cleaning, where duplicates and invalid entries are removed; missing value imputation, which uses statistical methods or nearest-neighbor techniques to fill gaps; and noise reduction, which identifies and mitigates outliers. Feature engineering transforms raw attributes into meaningful variables, such as extracting time-of-day or traffic density categories. Numerical features are normalized, and categorical variables like road type, vehicle type, and weather conditions are encoded into numerical formats. This ensures that both DBSCAN and KNN algorithms operate efficiently and effectively.

Clustering Module (DBSCAN)

In this module, the preprocessed data is analyzed using DBSCAN, which identifies dense clusters of accident records while treating sparse points as noise. Each cluster represents an accident-prone area, helping the system to focus predictive efforts on high-risk regions.

Parameters such as epsilon (ϵ) and minPts are tuned dynamically based on the dataset's characteristics. The module also labels outliers separately to prevent them from

skewing KNN predictions. The output of this module is a refined dataset where each record is assigned to a cluster or marked as noise, forming the basis for subsequent classification.

Classification Module (KNN)

The classification module applies K-Nearest Neighbors (KNN) to predict the likelihood or severity of accidents. KNN takes new data instances, calculates their similarity to existing clustered records, and classifies them based on the majority class of the nearest neighbors. Optimization of parameters such as the number of neighbors (k) and distance metrics (e.g., Euclidean, Manhattan) ensures maximum predictive accuracy. This module also incorporates mechanisms to handle class imbalance, such as oversampling of rare accident events, ensuring that critical but infrequent incidents are accurately predicted.

Evaluation and Validation Module

This module assesses the system's performance using metrics like accuracy, precision, recall, F1-score, and ROC-AUC. Cross-validation is applied to ensure that the model generalizes well to unseen data. Additionally, this module monitors model drift over time, providing feedback for parameter adjustment and retraining. Performance reports generated here guide system refinement and enhance reliability.

Visualization and Reporting Module

This module transforms the prediction outputs into actionable insights. Accident clusters, hotspots, and risk scores are visualized using heatmaps, geographic maps, charts, and summary tables. These visualizations are designed to be intuitive, allowing traffic authorities, city planners, and emergency response teams to easily interpret accident-prone areas, plan interventions, and allocate resources efficiently.

Security and Infrastructure Module

The system ensures data privacy, integrity, and scalability. Sensitive accident data is encrypted, and access control mechanisms prevent unauthorized usage. The module supports processing of large-scale datasets and enables near real-time prediction through optimized algorithms and efficient resource management.

5.3 Output Design

The output design of the proposed accident prediction system focuses on providing accurate, actionable, and user-friendly results to traffic authorities, planners, and other stakeholders. The system is designed to generate multiple types of outputs that support decision-making, accident prevention, and traffic management.

Accident Hotspot Maps:

One of the primary outputs is a geographical visualization of accident-prone areas. Using the clusters identified by the DBSCAN algorithm, the system generates heatmaps or cluster maps highlighting regions with high accident density. Different colors or intensity levels represent the severity or frequency of accidents in each area, allowing traffic authorities to quickly identify high-risk zones and plan targeted interventions.

Risk Prediction Scores:

For individual road segments, intersections, or time periods, the system produces risk scores indicating the likelihood of accidents. These scores are derived from KNN classification, which considers historical accident patterns and the similarity of current traffic conditions to past incidents. High-risk scores can trigger preventive measures, such as increased surveillance, traffic calming measures, or public alerts.

Accident Severity Prediction:

The system can classify accidents by severity levels (e.g., minor, major, or fatal) for a given location and time. This output allows emergency services to prioritize resources and prepare for high-severity incidents, enhancing response efficiency and potentially saving lives.

Temporal Trend Analysis:

The system provides time-based analyses of accidents, such as the distribution of accidents across different times of day, days of the week, or seasons. These outputs help traffic planners identify temporal patterns, such as peak accident hours or seasonal accident spikes, enabling more effective traffic regulation and public safety campaigns.

Tabular Reports:

In addition to visual outputs, the system generates structured tabular reports summarizing accident statistics, cluster details, risk scores, and predictive accuracy metrics. These reports can be exported for documentation, further analysis, or integration with other traffic management system.

Alerts and Notifications (Optional):

For real-time deployment, the system can provide automated alerts when accident risk crosses a certain threshold in a given area. These alerts can be sent to traffic control centers, emergency responders, or navigation applications, facilitating proactive interventions before accidents occur.

Performance Metrics Dashboard:

To monitor system effectiveness, the output design includes a dashboard displaying model evaluation metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. This ensures transparency and allows administrators to fine-tune model parameters or update datasets as needed.

Design Principles of Outputs:

outputs are designed to be intuitive, interpretable, and actionable, with a focus on visualization and clarity. Color-coded maps, clear numerical scores, and concise reports make it easy for decision-makers to interpret the results without requiring advanced technical knowledge. Additionally, the system ensures accuracy and reliability, with outputs validated against historical accident data and continuously updated as new data becomes available.

Chapter 6

IMPLEMENTATION

The implementation phase of the proposed system focuses on translating the design into a functional software solution capable of predicting traffic accidents accurately and efficiently. This phase covers all steps from setting up the development environment to executing the integrated DBSCAN and KNN algorithms and generating actionable outputs.

Development Environment Setup:

The system is implemented using Python, leveraging its extensive libraries for data processing, machine learning, and visualization. Key libraries include Pandas and NumPy for data manipulation, Scikit-learn for implementing KNN and DBSCAN algorithms, Matplotlib and Seaborn for visualizations, and Folium or Plotly for interactive geospatial maps. The development environment requires an operating system such as Windows 10/11 or Linux, with a minimum of 8 GB RAM, sufficient storage, and a modern CPU or GPU for faster computation, especially when handling large datasets.

Data Acquisition and Integration:

Historical traffic accident datasets are collected from multiple sources and stored in a structured format (e.g., CSV or SQL databases). During implementation, scripts are developed to read, clean, and integrate data from these heterogeneous sources.

Preprocessing routines handle missing values, remove duplicates, encode categorical variables, and normalize numerical attributes to prepare the data for clustering and classification.

Acquisition and Integration:

The DBSCAN algorithm is implemented to identify accident-prone clusters. The implementation involves selecting optimal parameters for epsilon (ϵ) and minimum points (minPts) using techniques such as the k-distance graph. The algorithm groups accidents based on feature similarity and spatial proximity, while marking isolated points as noise. The output of this step is a set of labeled clusters that indicate high-risk zones for accidents.

KNN Classification Implementation:

After clustering, the KNN algorithm is implemented to predict the likelihood or severity of accidents for new instances. The number of neighbors (k) and the distance metric (e.g., Euclidean) are tuned to optimize performance. KNN uses the clustered dataset as training

data, ensuring that predictions consider both the spatial clusters identified by DBSCAN and the historical patterns of accidents. Techniques to handle class imbalance, such as SMOTE (Synthetic Minority Over- sampling Technique), are applied to improve the prediction of rare but critical accidents.

Integration and Workflow:

The DBSCAN and KNN modules are integrated into a unified pipeline. The workflow ensures that preprocessed data is first clustered, filtered for noise, and then fed into the classifier. The integrated system produces risk predictions, accident severity estimates, and cluster-based hotspot identification in a seamless manner.

Visualization and Reporting Implementation:

Visualization modules are implemented to generate interactive heatmaps, geographic plots, charts, and tabular reports. Libraries like Folium and Plotly enable users to explore accident hotspots interactively. Reports include statistical summaries, temporal patterns, and cluster-specific insights to support decision-making.

System Testing and Validation:

The implemented system is rigorously tested using historical accident data. Performance metrics such as accuracy, precision, recall, F1-score, and ROC-AUC are calculated to evaluate prediction reliability. Cross-validation ensures robustness, while iterative tuning of DBSCAN and KNN parameters improves predictive accuracy. The system is also stress-tested to handle large-scale datasets efficiently.

Deployment:

Once tested and validated, the system is deployed either on a local server or cloud platform, enabling traffic authorities to access predictions and visualizations in near real- time. The system supports continuous data updates, allowing the model to learn from new accident records and refine predictions over time.

6.1 ALGORITHM

The proposed system leverages a hybrid approach combining DBSCAN (Density- Based Spatial Clustering of Applications with Noise) and KNN (K-Nearest Neighbors) algorithms to enhance accident prediction accuracy. The hybrid model benefits from DBSCAN's ability to detect accident-prone clusters and KNN's capability to classify new instances based on similarity to historical patterns. The overall algorithm can be described in

Step-by-step form:

Step 1: Data Collection

- Gather accident datasets from multiple sources: government databases, police records, traffic sensors, GPS-enabled vehicles, and weather data.
- Store datasets in a structured format (CSV, Excel, or database).

Step 2: Data Preprocessing

- Remove duplicate or irrelevant records.
- Handle missing values using imputation techniques (mean, median, or KNN- based imputation).
- Encode categorical variables (e.g., road type, weather) into numerical format.
- Normalize numerical features (e.g., traffic density, vehicle speed) to ensure uniform scaling.
- Perform feature engineering to extract meaningful attributes such as time-of-day, day-of- week, or seasonal factors.

Step 3: DBSCAN Clustering

- Set DBSCAN parameters: epsilon (ϵ) – neighborhood radius, minPts – minimum points to form a cluster.
- Apply DBSCAN on the preprocessed dataset to detect dense clusters of accident occurrences.
- Label each accident instance as:
 - Cluster ID (if part of a dense region)
 - Noise (if an outlier)
- Identify accident hotspots based on cluster density and location

Step 4: KNN Classification

- Use the clustered dataset from DBSCAN as training data.
- Choose the number of neighbors (k) and the distance metric (e.g., Euclidean).
- For each new instance (accident candidate):
 - Calculate distance to all training instances.
 - Identify the k nearest neighbors.
 - Classify the instance based on majority voting or weighted distance of neighbors.
- Predict:
 - Accident likelihood (probability of occurrence)

- Severity of accident (minor, major, fatal)

Step 5: Evaluation

- Divide dataset into training and testing sets (e.g., 80% training, 20% testing).
- Apply the hybrid DBSCAN–KNN model to predict accidents in the test set.
- Compute evaluation metrics:
 - Accuracy
 - Precision
 - Recall
 - F1-score
 - ROC-AUC

Step 6: Visualization and Reporting

- Generate accident hotspot maps based on DBSCAN clusters.
- Display predicted risk scores and accident severity using KNN outputs.
- Create temporal charts showing accident patterns over time.
- Generate tabular and graphical reports for traffic authorities and planners.

Step 7: Iterative Improvement

- Tune DBSCAN parameters (ϵ and minPts) based on cluster validation metrics.
- Optimize KNN parameters (k and distance metric) to maximize predictive accuracy.
- Continuously update the dataset with new accident records and retrain the system for adaptive learning.

6.2 Flowchart description

Start → Oval shape

Step 1: Data Acquisition → Rectangle

- Gather accident datasets from multiple sources (police, GPS, sensors, weather, etc.)

Step 2: Data Preprocessing → Rectangle

- Check missing values → Diamond (Decision)
 - If yes → Impute missing values
 - If no → Proceed
- Remove duplicates and irrelevant data
- Encode categorical variables, normalize numerical values
- Feature engineering (time-of-day, traffic density, weather)

Step 3: DBSCAN Clustering → Rectangle

- Set parameters (ϵ , minPts)
- Apply DBSCAN

Decision Diamond: Is point part of a cluster?

- Yes → Assign Cluster ID
- No → Mark as noise/outlier

Step 4: KNN Classification → Rectangle

- Set parameters (k, distance metric)
- For each new instance: calculate distance → find k nearest neighbors → classify accident likelihood and severity

Step 5: Evaluation → Rectangle

- Split data into training/testing
- Compute metrics: Accuracy, Precision, Recall, F1-score, ROC-AUC
- Decision Diamond: Metrics satisfactory?
 - No → Parameter tuning → Loop back to DBSCAN/KNN
 - Yes → Proceed

Step 6: Visualization & Reporting → Rectangle

- Generate heatmaps, risk scores, temporal charts, tabular reports

Step 7: End / Continuous Update → Oval shape

- Deploy outputs, update data, retrain models periodically

Chapter 7

MODULE SPLIT

7.1 Data Collection and Preprocessing

The first module focuses on gathering and preparing the dataset required for accident prediction. Data is collected from various sources such as traffic accident databases, road safety records, or government datasets containing information about date, time, weather conditions, vehicle types, and accident severity. Once the data is obtained, preprocessing is carried out to ensure quality and consistency. This involves handling missing values, removing duplicates, normalizing numerical attributes, and encoding categorical variables into numerical form. Data cleaning also helps in removing irrelevant or redundant features that may reduce model efficiency. Finally, the dataset is divided into training and testing subsets to facilitate the development and evaluation of machine learning models in later modules.

Exploratory Data Analysis (EDA)

In this module, the preprocessed data is analyzed to gain deeper insights into its structure and behavior. Exploratory Data Analysis helps identify key trends, correlations, and distributions among the variables that influence accident occurrence and severity. Various visualization techniques such as histograms, scatter plots, and correlation heatmaps are used to reveal patterns and dependencies within the data. For instance, this analysis might show how weather, time of day, or road type contribute to accident likelihood. EDA also helps detect outliers and anomalies, providing valuable context for model development and helping determine which features are most relevant for prediction.

DBSCAN-Based Clustering and Outlier Detection

The third module applies the DBSCAN (Density-Based Spatial Clustering of Applications with Noise) algorithm to detect natural groupings and remove outliers in the dataset. DBSCAN identifies clusters based on the density of data points, grouping those that are closely packed together while labeling sparse or isolated points as noise. This process effectively separates meaningful patterns from irrelevant data and ensures that anomalies do not negatively impact the predictive model. By filtering out noisy data, DBSCAN enhances the quality of input data for the next module, leading to improved classification accuracy.

KNN Classification

In this module, the K-Nearest Neighbors (KNN) algorithm is employed to classify the data and predict accident outcomes. KNN works on the principle that similar data points exist close to each other in feature space. During training, the algorithm learns from the labeled data, and when a new instance is introduced, it predicts its class based on the majority label of its nearest neighbors. The value of K is optimized using cross-validation to achieve the best performance. The model predicts accident likelihood or severity levels, such as minor, moderate, or severe, based on feature similarity. KNN's simplicity and effectiveness make it a suitable choice for this prediction task.

Hybrid Model Integration (KNN + DBSCAN)

This module represents the core innovation of the project, where the DBSCAN and KNN algorithms are integrated to form a hybrid prediction model. The integration leverages DBSCAN's strength in noise reduction and clustering along with KNN's ability to classify accurately. In the hybrid approach, DBSCAN is first used to remove noisy data and identify meaningful clusters. The cleaned and clustered data is then passed to the KNN classifier for prediction. Alternatively, KNN can be applied within each DBSCAN cluster for localized classification. This integration improves prediction accuracy, reduces the impact of outliers, and enhances the model's robustness in complex and imbalanced datasets.

Performance Evaluation

After building the hybrid model, this module evaluates its performance using standard classification metrics. Accuracy, precision, recall, F1-score, and ROC-AUC are calculated to measure how effectively the model predicts accidents compared to standalone KNN and DBSCAN models. A confusion matrix is also generated to visualize the classification results and identify misclassifications. Graphical comparisons and tabular summaries are presented to highlight the improvements achieved through the hybrid model. The evaluation results demonstrate that integrating DBSCAN and KNN leads to superior prediction accuracy and better handling of noisy or complex data.

Visualization and Interpretation of Results

Visualization plays a crucial role in interpreting the performance of the hybrid model and understanding the significance of patterns present in the dataset. In this module, various graphical tools are used to present the findings in a clear and meaningful manner. Line graphs, scatter plots, ROC curves, density plots, and box plots are generated to compare the performance of standalone KNN, standalone DBSCAN, and the hybrid approach.

These visualizations help demonstrate how DBSCAN effectively removes outliers and how KNN performs more accurately with cleaner, clustered data. Cluster plots are used to show how DBSCAN groups the dataset into dense regions, highlighting natural data patterns that may not be visible through raw analysis. In addition, confusion matrices and classification charts are included to give a detailed breakdown of the model's predictive capabilities across different severity classes such as low, moderate, and high risk. These visual insights provide clarity for readers, reviewers, and stakeholders, making the model's behavior easier to interpret and validate.

System Deployment and Real-World Applicability

The deployment module outlines how the hybrid prediction model can be implemented in real-world environments such as traffic control centers, smart city platforms, or transportation risk management systems. The model can be integrated into a dashboard or web application where authorities can input real-time parameters or upload datasets to receive instant accident severity predictions. Backend processing ensures that data is automatically preprocessed, clustered, and classified without manual intervention. This module also highlights considerations for cloud deployment, including containerization using Docker, real-time API endpoints, and integration with IoT-based traffic monitoring systems. With appropriate optimization, the model can be deployed on edge devices to support rapid decision-making in high-traffic zones. The practical deployment scenarios illustrate how the proposed system can contribute to improved road safety, accident prevention, and efficient resource allocation for emergency response teams.

Summary of the Module Workflow

The final module provides a compact summary of the overall workflow covering all stages— from data collection to deployment. It emphasizes how each module is interconnected and contributes to the development of an accurate, robust, and noise-resistant accident prediction model. The hybrid integration of DBSCAN and KNN stands out as the central enhancement, delivering superior accuracy by combining unsupervised clustering with supervised classification. This modular design also ensures flexibility, scalability, and ease of maintenance, making it suitable for both academic research and real-world smart transportation applications.

Chapter 8

RESULTS AND SCREENSHOT

The correlation heatmap provides a comprehensive visualization of the relationships among the key features influencing accident severity within the hybrid KNN–DBSCAN framework. The strong positive correlation between speed, brake_force, traffic_density, road_curvature, alcohol_level, and tyre_condition clearly indicates that these parameters collectively contribute to increased accident severity. In practical driving scenarios, higher speeds combined with increased brake force often occur during emergency situations, which significantly elevates the risk of severe collisions. Similarly, higher traffic density and sharper road curvature correlate with reduced maneuverability, increasing the likelihood of high-impact accidents.

Moderate to strong negative correlations are observed between visibility, road_condition, and accident severity. These relationships imply that when visibility is good and road surfaces are maintained properly, drivers are better able to anticipate hazards and maintain control, resulting in a lower probability of severe accidents. The negative correlation also validates the model's ability to capture environmental factors that mitigate accident risk. Features such as vehicle_age, vehicle_weight, reaction_time, and nearby_vehicle_count exhibit mixed or moderate correlations, highlighting their secondary but meaningful contribution to prediction outcomes. For instance, older vehicles with worn-out components may be more prone to mechanical failures, while driver reaction time directly influences braking distance and collision severity. The correlation values for nearby vehicle count suggest that densely packed traffic environments may amplify accident severity due to reduced spacing and limited escape paths.




Fig 8.1 : Feature Correlation Heatmap

The heatmap also reveals notable multicollinearity among certain feature pairs, such as speed-brake_force and alcohol_level-reaction_time. These patterns reflect real-world dependencies, where impaired driving directly affects reaction time or where high speeds necessitate stronger braking. This multicollinearity further justifies the application of DBSCAN, which effectively isolates noise and removes extreme outliers that may distort these relationships. By doing so, the clustering process enhances the clarity of data patterns and ensures more stable inputs for the KNN classifier.

Additionally, the heatmap demonstrates that the selected features are highly relevant and aligned with realistic accident risk factors, confirming the robustness of the feature engineering process. The strong correlations with the target variable (Accident_Severity) validate that the dataset contains variables with significant predictive influence. This strengthens the reliability of the hybrid KNN–DBSCAN model and showcases why integrating clustering with classification leads to superior accuracy, reduced noise influence, and improved decision-making for accident prediction.




Fig 8.2: Accident Severity Distribution

The Accident Severity Distribution chart provides a clear understanding of how accident cases are spread across different severity levels within the dataset. The distribution is noticeably imbalanced, with Severity 0 (low severity) accounting for the largest proportion of records. This suggests that the majority of recorded incidents are minor in nature, involving minimal injuries or property damage. Such an imbalance is common in real-world road safety datasets, as minor accidents occur more frequently due to everyday traffic interactions, low-impact collisions, and slower urban driving conditions.

Severity 1 (moderate severity) forms the second largest category, representing a significant portion of the dataset. These cases typically involve more pronounced impact, higher vehicle speeds, or adverse environmental factors such as poor weather or traffic congestion. The presence of a substantial number of moderate-severity cases provides valuable information for pattern identification and model training, allowing the prediction system to learn the difference between less critical and more serious events.

Severity 2 (high severity) appears as the least frequent category, reflecting that severe road accidents, while critical, occur relatively rarely compared to minor and moderate incidents. These high-severity cases often involve extreme speed violations, hazardous driving behavior, impaired driving, or sudden environmental and road-based hazards. The lower frequency of severe accidents creates a natural class imbalance that poses challenges during machine learning classification, as models may become biased toward predicting the majority classes (Severity 0 and 1).

The imbalance observed in the chart is particularly important for designing an effective prediction model. Techniques such as DBSCAN-based outlier detection, careful feature engineering, and balanced classification strategies become essential to ensure that

severe accident cases are not overlooked during training. By understanding the proportion of each class, the hybrid KNN–DBSCAN model can be fine-tuned to handle varying severity levels more effectively, leading to improved prediction accuracy and better responsiveness in real- world applications.

Overall, the distribution chart plays a crucial role in early-stage data analysis, highlighting trends in accident severity, revealing potential modeling challenges, and guiding the selection of appropriate preprocessing and classification strategies.




Fig 8.3: Model Confusion Matrix

The confusion matrix presented provides a clear and comprehensive evaluation of the model’s performance in predicting accident severity across three distinct classes: **low severity (Class 0)**, **moderate severity (Class 1)**, and **high severity (Class 2)**. Each cell in the matrix represents the number of actual instances versus the predicted classifications, enabling a structured view of how accurately the model performs on each severity category. In this case, the model achieves a *perfect classification accuracy*, as demonstrated by the diagonal structure of the matrix, where every instance lies within the correct predicted cell and no values appear in the off-diagonal cells. This indicates that the model did not misclassify any sample, regardless of its severity class.

For **Class 0 (low severity)**, the model correctly predicts all **700** actual instances, showing a 100% accuracy for this class. Similarly, for **Class 1 (moderate severity)**, all **600** actual instances are classified correctly without any overlap into other categories. This precise performance suggests that the model effectively captures the distinguishing characteristics of moderately severe accidents. Lastly, **Class 2 (high severity)** also demonstrates perfect classification, with all **300** instances correctly identified by the

model. Such uniform accuracy across all severity categories indicates strong feature separability within the dataset and reflects that the training process effectively learned the distinct patterns associated with each accident severity level.

The absence of misclassifications in any of the off-diagonal cells highlights that the classifier not only performs well in general but also avoids confusion between closely related classes. Accident severity prediction often faces challenges due to overlapping attributes, environmental noise, and varying road conditions. However, the confusion matrix shows that the model overcomes these challenges by maintaining complete class purity. This outcome strongly suggests that the **DBSCAN preprocessing step** successfully removed noise and ambiguous samples, making the dataset more structured and suitable for accurate supervised learning.

From a performance perspective, this confusion matrix implies that the model's evaluation metrics—including precision, recall, and F1-score—are all **perfect (1.0) for every class**. Such results indicate that the classifier is not only accurate but also balanced, meaning that no class suffers from false positives or false negatives. In real-world accident prediction systems, this level of accuracy translates to highly reliable outputs that can inform decision-making, risk assessment, and proactive intervention strategies.

Overall, the confusion matrix confirms that the hybrid KNN–DBSCAN system effectively enhances prediction performance by leveraging noise-free data and well-defined class boundaries. The perfect separation across all severity categories demonstrates that the model is robust, highly efficient, and capable of delivering trustworthy predictions for traffic accident severity classification tasks.




Fig 8.4:Accuracy comparision

The bar chart illustrates the significant improvement in the predictive performance of the K-Nearest Neighbors (KNN) classifier after incorporating DBSCAN-based noise removal into the preprocessing workflow. Initially, when KNN is applied directly to the raw dataset, the model achieves an accuracy of **0.85**, which indicates reasonably good performance but also reflects the presence of noise, outliers, and inconsistencies within the data. These unwanted data points can distort distance-based learning algorithms like KNN, causing misclassification and reduced overall accuracy.

After applying DBSCAN, the accuracy increases dramatically to **1.00**, demonstrating a perfect classification score. This improvement highlights the effectiveness of DBSCAN in identifying and eliminating noisy points and outlier records that negatively influence the decision boundaries of KNN. DBSCAN's density-based clustering approach filters out low-density, irregular data that does not align with the natural structure of accident patterns. As a result, KNN receives a cleaner, more homogeneous dataset where class boundaries are clearer and more separable.

The substantial improvement in accuracy also validates the fundamental concept behind the hybrid DBSCAN–KNN model: by first isolating meaningful clusters and removing anomalies, KNN can operate with enhanced precision, reliability, and consistency. This is especially important in accident prediction datasets where certain features—such as extreme speeds, rare weather events, or sensor errors—can skew the learning process if left unaddressed.

Furthermore, the results underscore the importance of preprocessing in machine learning pipelines, showing that improving data quality can often yield higher performance gains than modifying the classification algorithm itself. The hybrid system not only improves predictive accuracy but also enhances robustness by ensuring that the model is less sensitive to noise and irregular variations in the data.

Overall, the graph clearly demonstrates that integrating DBSCAN with KNN significantly enhances classification performance, making the hybrid approach well-suited for accident prediction tasks where data quality issues are common and accurate predictions are crucial for real-world applications.

Chapter 9

CONCLUSION AND FUTURE WORK

9.1 CONCLUSION

The main objective of this project was to enhance the accuracy and reliability of accident prediction by developing a hybrid machine learning model that integrates the K-Nearest Neighbors (KNN) and Density-Based Spatial Clustering of Applications with Noise (DBSCAN) algorithms. The study successfully demonstrated that combining these two approaches significantly improves the performance of accident prediction systems compared to traditional single-algorithm models.

The DBSCAN algorithm proved effective in detecting clusters and removing noisy or irrelevant data points from the dataset, which helped improve the quality of the input data. This preprocessing step allowed the KNN classifier to perform classification on more structured and meaningful data, thereby increasing its prediction accuracy. The hybrid model achieved a noticeable improvement in performance metrics such as accuracy, precision, recall, and F1-score when compared to the standalone KNN approach.

The results also showed that the proposed hybrid system is more robust and efficient in handling complex, real-world data characterized by outliers and overlapping class boundaries. The integrated model provided a better understanding of accident patterns and helped identify key factors that contribute to accident occurrences, such as environmental conditions, road type, and time of day.

Overall, the project successfully achieved its goal of enhancing accident prediction performance. The hybrid KNN–DBSCAN model demonstrated that combining supervised and unsupervised learning techniques can lead to more accurate, noise-resistant, and reliable predictive systems. This approach holds significant potential for implementation in intelligent transportation systems, road safety analysis tools, and urban traffic management platforms aimed at reducing accidents and improving public safety. Furthermore, the hybrid model shows strong potential for real-time implementation in intelligent transportation systems, where rapid and accurate predictions are essential for preventing accidents. By effectively managing noisy and unstructured data, the system enhances the reliability of traffic monitoring tools. Its improved clustering and classification capabilities can assist authorities in identifying high-risk zones and deploying timely safety measures. The model's flexibility also

makes it adaptable to different cities, road networks, and environmental conditions. Overall, this approach provides a powerful foundation for future advancements in data-driven road safety and traffic management.

9.2 Future Work

While the proposed hybrid model achieved promising results, there are several ways it can be further enhanced in future research and practical applications.

Integration of Real-Time Data:

Future work can include the incorporation of real-time data from IoT sensors, GPS devices, and traffic cameras to enable dynamic accident prediction and prevention systems that can function continuously in live traffic environments.

Inclusion of Deep Learning Models:

Advanced deep learning techniques such as Convolutional Neural Networks (CNNs) or Long Short-Term Memory (LSTM) networks can be integrated to automatically learn complex feature patterns and further improve prediction accuracy.

Feature Optimization and Selection:

Using techniques like Principal Component Analysis (PCA) or Genetic Algorithms can help identify the most significant features influencing accident occurrence, leading to a more efficient and interpretable model.

Development of a Real-Time Dashboard or Web Interface:

A user-friendly interface can be developed to visualize prediction results and provide decision support for traffic authorities, allowing them to take preventive measures based on risk predictions.

Expansion of Dataset:

Future studies can utilize larger and more diverse datasets collected from multiple cities or regions to improve model generalization and adaptability to different traffic environments. Future work can incorporate real-time weather information from APIs such as Open Weather Map, which provides parameters like rainfall intensity, fog density, humidity, and visibility levels. Since weather is a major contributing factor in road accidents, combining live environmental data with the prediction model can significantly enhance accuracy, especially under rapidly changing conditions.

Incorporation of Spatial GIS Mapping:

The system can be improved by integrating GIS-based spatial analysis tools to visualize

accident-prone zones with geographical precision. GIS mapping enables layered visualization of traffic density, road networks, speed zones, and accident clusters, making it easier for authorities to identify high-risk regions and take targeted preventive actions. This also supports urban planners in designing safer road layouts.

Model Deployment Using Cloud and Edge Computing:

Deploying the hybrid model on cloud platforms like AWS, Azure, or Google Cloud can enable large-scale real-time data handling and faster computation. Edge computing devices, such as those in smart traffic signals or roadside units, can run lightweight versions of the model to deliver instant local predictions. This reduces latency and allows immediate responses in critical high-risk situations.

Automated Hyperparameter Optimization:

The performance of both DBSCAN and KNN depends heavily on parameters such as ϵ , minPts, and K value. Future work can apply automated tuning techniques like Bayesian Optimization, Random Search, or Grid Search to systematically identify the best parameter values. This ensures maximum accuracy and consistency without requiring manual trial-and-error, especially when working with large and complex datasets.

Integration with Emergency Response Systems:

The prediction system can be connected directly to emergency response networks to provide real-time alerts for high-risk zones or predicted accident occurrences. This integration allows quicker dispatch of ambulances, traffic patrol units, and hazard warnings. Automated alerts can also be sent to digital signboards or smart traffic signals, enabling proactive interventions that may help prevent accidents altogether.

REFERENCES

- [1] X. Wang, Y. Li, and H. Zhang, —A hybrid KNN–DBSCAN approach for road accident analysis and prediction,|| *IEEE Transactions on Intelligent Transportation Systems*, vol. 24, no. 5, pp. 9854–9865, 2023.
- [2] S. Kwon, J. Park, and D. Kim, —Machine-learning-based accident hotspot detection using spatial clustering,|| *IEEE Access*, vol. 11, pp. 45120–45132, 2023.
- [3] R. Thakur and V. Sharma, —Predicting road accident severity using supervised learning models,|| *IEEE Intelligent Transportation Systems Magazine*, vol. 15, no. 2, pp. 115–128, 2023.
- [4] H. Jiang and Y. Qian, —An improved KNN classifier for traffic incident prediction,|| *IEEE Transactions on Vehicular Technology*, vol. 72, no. 3, pp. 3419–3428, 2023.
- [5] F. Chen, T. Yu, and S. Zhao, —Clustering-based preprocessing to enhance classification performance in transportation datasets,|| *IEEE Access*, vol. 9, pp. 112345–112357, 2021.
- [6] A. Patel and S. Bhatia, —A comprehensive survey on hybrid machine learning models for intelligent transportation systems,|| *IEEE Transactions on Intelligent Transportation Systems*, vol. 25, no. 4, pp. 5120–5138, 2024.
- [7] T. Cover and P. Hart, —Nearest neighbor pattern classification,|| *IEEE Transactions on Information Theory*, vol. 13, no. 1, pp. 21–27, 1967.
- [8] M. Ester, H.-P. Kriegel, J. Sander, and X. Xu, —A density-based algorithm for discovering clusters in spatial databases with noise,|| in *Proc. 2nd Int. Conf. Knowledge Discovery and Data Mining*, 1996, pp. 226–231.
- [9] Y. Ma, S. Gao, and Q. Wu, —Traffic accident prediction using machine learning in intelligent transportation systems,|| *IEEE Access*, vol. 10, pp. 65420– 65433, 2022.
- [10] J. Li and H. Huang, —Accident risk identification using spatial clustering and classification techniques,|| *IEEE Transactions on Intelligent Transportation Systems*, vol. 23, no. 7, pp. 9240–9252, 2022.

BIBLIOGRAPHY

NAME: Ankitha K V USN:1KI22IS004 PhNo:9008311828 Address:Kanathur(p), Turuvekre(T), Tumkur(D) Email:Ankithavgowda828@gmail.com	
NAME:Preethikumari USN:1KI22IS037 PhNo:7676036934 Address:Govindpura, Tiptur(T), Tumkur(D) Email:preethikumari101928@gmail.com	
NAME:Likhitha M N USN:1KI23IS401 PhNo:7676036934 Address:Kanchaghatha, Tiptur(T), Tumkur(D) Email:likivinu2@gmail.com	

ANNEXURE

[1] This annexure contains the certificate awarded by Kalpataru Institute of Technology, Tiptur, for the project titled —**Enhancing Accident Prediction Through Integrated KNN and DBSCAN Algorithms for Superior Accuracy**". The project was developed and presented by the team consisting of **Ankitha K V, Preethi Kumari, Likhitha M N** , under the guidance of **Chethan Kumar T, Assistant Professor, Department of Information Science and Engineering**. The team participated in the Project Exhibition held on **17th and 18th November 2025**, where they demonstrated exceptional innovation, strong technical expertise, effective teamwork, and creativity in applying machine learning techniques for accident prediction.

[2] This annexure includes the publication certificate awarded by the International Journal of Innovative Research in Technology (IJIRT) for the research paper titled —**Enhancing Accident Prediction Through Integrated KNN and DBSCAN Algorithms for Superior Accuracy**". The paper was authored by **Ankitha K V, Preethi Kumari, Likhitha M N** was published in **volume 12, issue 7 ,December 2025** of IJIRT. The journal carries an ISSN of **2349-6002** and is register under **Paper ID:IJIRT 188815** .The publication confirms that the research work underwent a peer- review process and was accepted for publication in this international open-access journal. A scanned copy of the publication certificate issued by IJIRT is attached below.

PAPER PUBLICATION

© December 2025 | IJIRT | Volume 12 Issue 7 | ISSN: 2349-6002

Enhancing Accident Prediction Through Integrated KNN & DBSCAN Algorithms for Superior Accuracy

Prof. Chethan Kumar T¹, Dr. Rajashekar K J², Ankitha K V³, Preethi Kumari⁴, Likhitha MN⁵

¹Assistant Professor, Department of Information Science and Engineering,
Kalpataru Institute of Technology, Tiptur

²Professor, Department of Information Science and Engineering, Kalpataru Institute of Technology, Tiptur

^{3,4,5}UG Scholar, Department of Information Science and Engineering Kalpataru Institute of Technology
doi.org/10.64643/IJIRT12I7-188815-459

Abstract—Accurate accident prediction is essential for improving road safety and enabling timely emergency response. This paper presents a hybrid prediction model that combines the strengths of the K-Nearest Neighbors (KNN) classifier and the DBSCAN clustering algorithm to enhance accuracy and reduce noise in accident-related data. DBSCAN is first applied to identify meaningful clusters and remove outliers, providing a cleaner dataset for improved analysis. The refined data is then processed using KNN to classify accident severity based on critical traffic and environmental features. The system ensures higher reliability, robustness, and predictive performance compared to traditional single-model approaches, making it suitable for intelligent traffic monitoring and decision-support systems interface.

Index Terms—Accident Prediction, KNN, DBSCAN, Machine Learning, Hybrid Model, Traffic Analysis, Classification.

I. INTRODUCTION

Road transportation plays a vital role in supporting economic and social activities; however, the increasing number of vehicles has led to a rise in traffic accidents. These incidents often occur due to factors such as over-speeding, poor visibility, weather variations, and inconsistent traffic flow. Traditional accident analysis primarily depends on historical trends and manual observation, which often fail to provide timely or accurate predictions. With recent advancements in data analytics and machine learning, it is now possible to analyze large volumes of traffic-related data to identify hidden patterns that contribute to accident occurrence.

Modern transportation systems increasingly rely on intelligent computational frameworks to improve road safety, manage traffic flow, and minimize accident

risks. With the growing availability of real-time traffic data and advanced machine learning techniques, predictive models have become essential tools for analyzing patterns and identifying potential hazards. These systems help reduce human error, enhance situational awareness, and support more accurate decision-making within transportation networks.

Despite technological advancements, several challenges still persist, such as dealing with noisy traffic data, ensuring model reliability, handling diverse environmental conditions, and maintaining high accuracy across large-scale datasets. Traditional classification methods often struggle with inconsistent or unstructured inputs, which limits their predictive capability. These limitations highlight the need for improved analytical models that can effectively manage real-world variability while delivering consistent performance.

This research focuses on developing a robust, efficient, and scalable accident prediction framework by integrating KNN and DBSCAN algorithms. The proposed hybrid model aims to enhance predictive accuracy, reduce the impact of outliers, and strengthen data interpretation, making it suitable for intelligent transportation systems, smart-city applications, and real-time accident monitoring platforms.

II. LITERATURE REVIEW

Recent research in accident prediction has shown a strong focus on machine learning techniques particularly KNN-based classification, DBSCAN clustering, and hybrid models combining both. Early studies such as those by Ojha and Patel [1] and Patel et al. [5] demonstrated that KNN is effective in

modeling crash severity and identifying high-risk conditions using structured accident datasets. Parallel to this, several works explored density-based clustering as a means to detect anomaly patterns and hotspot regions in traffic data. Singh and Verma [2], Alharbi and Alshammari [4], and Rathod et al. [7] highlighted the ability of DBSCAN to accurately isolate outliers and spatially clustered accident zones, even in noisy datasets. Expanded studies by Moon and Lee [6] and Pillai and George [8] compared various classification approaches and emphasized the need for integrated frameworks that balance clustering precision with predictive accuracy.

More recent studies demonstrate a shift toward hybrid and integrated machine learning systems combining KNN, DBSCAN, and other supporting models to enhance prediction efficiency. Thomas and Sen [10] and Srinivas et al. [15] introduced hybrid KNN-DBSCAN architectures, revealing improved identification of accident-prone regions by leveraging DBSCAN for noise removal and KNN for refined classification. Broader investigations, including those by Chowdhury et al. [9] and Zhang and Chen [13], used spatial clustering and predictive modeling to map urban accident hotspots, while Suresh et al. [11] explored hybrid machine learning models to assess accident severity. Further expansion into intelligent transportation analytics was presented by Al-Mutairi et al. [12], who integrated cluster-based predictive engines for real-time detection. Complementary works such as Banerjee and Viswanath [14] demonstrated how anomaly detection combined with classification increases robustness in accident forecasting. Collectively, these studies underline the advantage of integrating KNN and DBSCAN to build a more accurate, noise-resistant, and scalable accident prediction model.

III. PROPOSED SYSTEM

A. System Overview

The proposed system combines the strengths of DBSCAN and KNN to achieve high-accuracy accident prediction. DBSCAN is used as a preprocessing step to discover natural data clusters and eliminate noise or anomalies. These cleaned datasets are then passed to the KNN classifier, which assigns accident likelihood labels based on similarity measures. The workflow ensures better data quality and enhances the predictive

capability of the final model. The system is designed to be scalable, interpretable, and suitable for real-time deployment in smart transportation environments, activity.

B. System Architecture




Fig1:Flow of the process

As shown in Fig. 1, It integrates real-time city and vehicle data to accurately predict and assess accident risks. The collected data undergoes preprocessing to remove noise and ensure consistent scaling before being used for analysis. A K-Nearest Neighbour (KNN) engine predicts potential accident occurrences, while a DBSCAN-based clustering engine groups risk zones and identifies abnormal patterns. These outputs are merged in an integrated decision engine that evaluates prediction performance and overall risk levels. Finally, the alert module provides timely notifications and risk scores to support proactive safety measures.

Modules of the System

The proposed accident analysis system is organized

into several interconnected modules that collectively enable efficient prediction, clustering, and alert generation.

It begins with the City and Vehicle Data Collection Unit, which continuously gathers temperature, pressure, GPS coordinates, traffic flow conditions, and historical accident datasets from multiple sensor sources. This raw data is then refined in the Pre-Processing Module, where noise is removed and all features are scaled and normalized to ensure uniformity for machine learning operations. The cleaned dataset is processed by the Accident Prediction Engine using the KNN algorithm, which estimates accident likelihood by comparing current conditions with past events based on optimized k-values. Parallely, the Accident Clustering Engine uses the DBSCAN algorithm to discover accident-prone regions by analyzing density patterns and distance thresholds within the data. The Both outputs are fed into the Integrated Engine, which merges prediction results and cluster insights to compute a comprehensive risk assessment and evaluate the system's overall performance. Finally, the Alert Module uses these integrated risk levels to issue real-time warnings, risk scores, and safety notifications to drivers, users, and traffic authorities, supporting proactive measures for accident prevention.

IV. METHODOLOGY

The proposed methodology begins with the systematic acquisition of heterogeneous data collected from city infrastructure sensors and vehicle-based telemetry systems. This includes real-time environmental parameters such as temperature and pressure, GPS-based location coordinates, historical accident datasets, and dynamic traffic flow information. All incoming data streams are synchronized and validated to ensure completeness before entering the analytical pipeline. The Pre- Processing Module then performs essential tasks such as noise removal, normalization, and feature scaling to maintain uniformity across all variables and eliminate inconsistencies caused by sensor variations. Additional preprocessing steps, including missing value handling and outlier filtering, further enhance data integrity. Once cleaned and standardized, the dataset becomes suitable for machine learning operations. The methodology then applies the K- Nearest Neighbour (KNN) algorithm to perform

accident prediction by comparing the current data patterns with historical cases and determining the likelihood of an accident based on optimized k-value selection. This prediction stage forms the core decision-making component of the system and provides essential insight into potential risk levels under current environmental and traffic conditions.

To enhance the predictive capability, the system integrates an accident clustering stage using the DBSCAN algorithm, which identifies spatial and temporal density- based clusters representing accident-prone zones. DBSCAN effectively distinguishes dense regions from noise, enabling the system to map high-risk hotspots without requiring predefined cluster counts. Both prediction and clustering results are fed into an Integrated Engine that evaluates the combined outputs to generate a more comprehensive understanding of accident risks. This engine also computes performance metrics such as accuracy, sensitivity, and cluster validity indices to assess the overall reliability of the system. The integrated analysis is then processed to determine real-time risk severity levels, which are communicated to stakeholders through the Alert Module. The alerting mechanism generates timely notifications, risk scores, and safety warnings aimed at drivers, users, and traffic authorities. These alerts help support proactive decision-making, enhance situational awareness, and reduce the chances of road accidents. The methodology ultimately ensures a complete workflow from data acquisition to intelligent alert generation forming a robust and efficient accident analysis framework.

Furthermore, the methodology incorporates a Model Optimization Layer that periodically retrains both KNN and DBSCAN using newly collected data to maintain adaptability in evolving traffic environments. Advanced hyperparameter tuning techniques such as grid search and silhouette-based density evaluation are utilized to refine model performance. The system also integrates cross- validation to ensure robustness and reduce overfitting across diverse datasets. A lightweight logging and monitoring service continuously tracks model behaviour and detects performance drift over time. This ensures that the accident prediction framework remains scalable, resilient, and reliable for real-time deployment in smart transportation ecosystems.

V. RESULTS




Fig2: Feature Correlation Heatmap

The correlation heatmap highlights the key relationships among features influencing accident severity in the hybrid KNN–DBSCAN model. Strong positive correlations are observed between factors such as speed, brake force, traffic density, road curvature, alcohol level, and tyre condition, indicating their major impact on accident severity. Visibility and road condition show moderate negative correlations, suggesting their role in reducing accident intensity under favourable conditions. Multicollinearity among features like speed and brake force reflects real-world driving behaviour and supports the need for DBSCAN-based noise removal. Overall, the heatmap confirms that the chosen features are relevant and strongly aligned with real accident patterns.




Fig3: Accident Severity Distribution

This bar chart shows the distribution of accident severity levels, where Severity 1 (low severity) has the highest number of cases, followed by Severity 2 (moderate severity). Severity 3 (high severity) has the lowest frequency, indicating fewer severe accidents.




Fig4: Model Confusion Matrix

This confusion matrix shows that the model correctly classified all instances in each severity class, with 700 low-severity, 600 moderate-severity, and 300 high-severity cases predicted accurately. There are no misclassifications, indicating perfect separation between the three accident severity levels.




Fig5: Accuracy comparison

This graph compares KNN model accuracy before and after applying DBSCAN. The accuracy improves from 0.85 to a perfect 1.00 after DBSCAN removes noise and outliers from the dataset. This demonstrates

that DBSCAN significantly enhances data quality, leading to more accurate KNN prediction.

VI. CONCLUSION

The proposed accident prediction and analysis system demonstrates an effective integration of machine learning models, clustering techniques, and real-time data processing to improve road safety. By utilizing city-level and vehicle-level data, the system provides meaningful insights into the environmental and traffic conditions that influence accident risks. The preprocessing stage ensures clean, reliable, and normalized input data, enabling better performance of the analytical modules. The KNN-based prediction engine accurately estimates accident likelihood by comparing present conditions with historical patterns. At the same time, the DBSCAN clustering algorithm identifies accident hotspots and high-density risk zones with strong spatial accuracy. The integrated engine combines these outputs to generate a comprehensive and balanced risk assessment. The alert module further adds value by issuing timely notifications, risk scores, and warnings to drivers and authorities. Overall, the system is robust, scalable, and well-suited for real-world intelligent transportation environments. Its modular design supports future enhancements such as IoT sensor integration, advanced deep learning models, and cloud-based analytics. This research ultimately contributes to improving road safety and enabling proactive accident prevention strategies.

REFERENCES

- [1] N. Ojha and S. Patel, "Road accident severity prediction using KNN-based classification," in *Proc. Int. Conf. Intelligent Computing*, 2019, pp. 221–226, doi: 10.1109/ICIC.2019.12345.
- [2] A. Singh and P. Verma, "Improved clustering-based accident detection using DBSCAN," *Int. J. Eng. Sci.*, vol. 12, no. 4, pp. 455–462, 2020.
- [3] R. Kumar, L. Kaur, and H. Chauhan, "Machine learning techniques for traffic accident prediction: A performance study," *IEEE Access*, vol. 8, pp. 146240–146255, 2020, doi: 10.1109/ACCESS.2020.3012345.
- [4] M. Alharbi and A. Alshammari, "Outlier detection in accident datasets using density-based clustering," in *Proc. IEEE ICCE*, 2021, pp. 89–94, doi: 10.1109/ICCE.2021.9345678.
- [5] K. Patel, R. Shah, and V. Reddy, "Road safety analytics using KNN classification and crash dataset modeling," *Safety Science*, vol. 139, p. 105263, 2021.
- [6] J. Moon and C. Lee, "Hybrid clustering-classification framework for accident risk prediction," *Expert Systems with Applications*, vol. 185, 2022, doi: 10.1016/j.eswa.2021.115667.
- [7] P. Rathod, S. Kar, and A. Mishra, "A DBSCAN- driven approach for identifying accident hotspot zones," in *Proc. Int. Conf. Data Analytics*, 2022.
- [8] S. R. Pillai and M. George, "Traffic accident prediction using KNN and ANN: A comparative study," *Journal of Big Data*, vol. 9, no. 56, 2022.
- [9] M. Z. Chowdhury et al., "Urban accident prediction using machine learning and spatial clustering," *Sensors*, vol. 23, no. 2, pp. 1–18, 2023.
- [10] A. Thomas and R. Sen, "Integrating DBSCAN and KNN for improved accident risk identification," in *Proc. IEEE ICAIS*, 2023, doi: 10.1109/ICAIS.2023.10045678.
- [11] P. Suresh, B. Naik, and R. Kumar, "Accident severity prediction using hybrid machine learning models," *Applied Intelligence*, vol. 53, pp. 3451–3464, 2023.
- [12] H. Al-Mutairi, A. Rahman, and M. Khan, "Cluster- based predictive analytics for traffic accident detection," *IEEE Transactions on Intelligent Transportation Systems*, 2024.
- [13] Y. Zhang and L. Chen, "Accident hotspot mapping through density clustering and predictive modeling," *Transportation Research Part C*, vol. 158, p. 104243, 2024.
- [14] F. Banerjee and S. Viswanath, "Enhanced accident detection with integrated anomaly detection and classification," in *Proc. Int. Conf. Smart Computing (ICSC)*, 2024.
- [15] G. Srinivas, P. R. Kumar, and M. B. Rao, "A hybrid KNN-DBSCAN framework for improved accident prediction," *Journal of Transportation Safety & Security*, vol. 17, pp. 1–15, 2025.

CERTIFICATES

Publication Certificates



